

Time-Domain Imaging of Rotating, Cloud-Covered Exoplanets with the Solar Gravitational Lens: Reconstruction Algorithms, Observing-Cadence Requirements, and Mission Targets

Sanskar Sontakke

Independent Researcher

sanskarsontakke@gmail.com ORCID: 0009-0008-4531-5707

ABSTRACT

The solar gravitational lens (SGL) can, in principle, deliver multipixel images of terrestrial exoplanets from heliocentric distances $z \simeq 650\text{--}900$ AU. In practice the planet rotates and its cloud cover evolves while a spacecraft raster-scans the kilometer-scale image cylinder one “pixel” at a time, so the measurements are temporally aliased; recent studies identify this dynamic inversion problem as a leading unsolved challenge for the concept. We present an open, end-to-end simulation framework that reproduces the published SGL photometric benchmarks (the aperture-averaged $d/4p$ kernel; per-sample $\text{SNR}_C = 43.16$ for 1800 s dwells on an exo-Earth at 30 pc from 650 AU; the $0.891 D/(d\sqrt{N})$ deconvolution penalty) and then treats rotation, orbital illumination, and stochastic, advecting clouds explicitly. We formulate image recovery as a regularized generalized least-squares inversion of the time-tagged sample stream (time-domain inversion, TDI, distinguished here from detector-level time-delay integration), in which rotation and illumination are modeled directly by the forward operator and cloud variability enters statistically. **For a rotating, cloud-free planet observed under known spin geometry, TDI recovers the surface albedo map with Pearson $r = 0.9786 \pm 0.0002$ (SSIM = 0.848 ± 0.001) over 10 paired seeds. Rotational aliasing is tractable under characterized ephemerides; it is not an information ceiling, and the correlation lost within the supported harmonic range is regularization, not spin. Clouds, not coronal photon noise, dominate the error budget: per-sample cloud-induced variance exceeds coronal shot noise by 163–440 (12.8–21.0 in standard deviation) at $t_s = 1800$ s, collapsing fidelity to $r = 0.249 \pm 0.010$ at 71% realized cover under single-campaign budgets. The observing-cadence law is measured under two explicit, mutually exclusive resource ledgers (fixed 8 h of photons per raster position, or fixed 89.87 d of wall clock): replacing 4 long dwells with 64 short revisits improves r from 0.049 to 0.571 at fixed photons, monotonically and unanimously across six paired seeds, while the two arms agree to within the paired standard error. Slew, settling and metrology overheads (45 s per dwell) reduce the duty cycle from 99.4% to 90.5% over that range. Ablation of the estimator itself shows that the algorithmic gain comes from eliminating the per-dwell-slot cloud nuisance: exact anchored profiling of the slot means lifts r by $+0.054 \pm 0.006$ over a fit that ignores them ($p = 1.1 \times 10^{-5}$, 10/10 seeds), whereas temporal whitening contributes nothing measurable at the fiducial cadence (-0.003 ± 0.002 , $p = 0.13$). The spin-ephemeris requirement is measured rather than asserted: a 5° pole tilt costs 0.057 in r and a 5° initial-phase error 0.047, and 60° of phase error destroys the map entirely. We withdraw the v2 claim that residual misalignments are refinable in flight via profile likelihood: the whitened χ^2 is exactly invariant to assumed initial phase and monotone in assumed period over the grid tested, so the spin state is observable through image quality but not recoverable from photometric residuals by that route. Surface correlation $r(\ell)$, computed with a real spherical-harmonic transform and $\cos \vartheta$ quadrature weights, stays above 0.99 for all degrees up to the $\ell = 20$ ceiling of the 36×72 grid when cloud-free, and under clouds falls below 0.5 at $\ell \simeq 9$ in a five-point running mean (~ 2200 km for an Earth-radius planet). A land–ocean separation of $d' = 0.73 \pm 0.03$ at 55.8% realized cover exceeds every member of a 40-sample physical cloud-only null, giving $p \leq 0.024$ — the ensemble floor, not a measurement of 0.02. Finally, we map requirements onto missions: we list the five nearest confirmed potentially habitable planets around M dwarfs plus τ Cet e as a reproducible nominal-proxy table (not an uncertainty propagation), evaluate image-plane dynamics ($0.11\text{--}5.8$ km s $^{-1}$ of along-track velocity change per 90 d), discuss deep-perihelion sailcraft, planar radioisotope tiles and optical relays, and swarms as an asymptotic architectural scenario.**

Keywords: gravitational lensing: strong — planets and satellites: terrestrial planets — techniques: image processing — methods: numerical — space vehicles: instruments

1 Introduction

Resolving the surface of an Earth-like exoplanet is beyond any credible filled-aperture or interferometric architecture: a 10×10 surface map of an exo-Earth at 10 pc requires resolution elements of ~ 0.85 μas , i.e., a $\sim 10^2$ km optical aperture, and integration times for conventional concepts reach $10^4\text{--}10^5$ yr once

realistic backgrounds are included (Turyshev 2026a). The solar gravitational lens (SGL) evades this scaling by using the Sun as the primary optical element (Eshleman 1979; Turyshev & Toth 2017). Wave-optical treatments show that a strong-interference region forms beyond $z_0 \simeq b^2/2r_g \simeq 547.8$ AU (for rays of impact parameter $b \simeq R_\odot$, with $r_g = 2GM_\odot/c^2 \simeq 2.95$ km), with

on-axis amplification $\mu_0 = 4\pi^2 r_g / \lambda \simeq 1.2 \times 10^{11}$ at $\lambda = 1 \mu\text{m}$ and angular resolution of order 10^{-10} arcsec (Turyshev & Toth 2017, 2019). An Earth-radius planet at 30 pc is compressed into an image cylinder only $D_{\text{img}} \simeq 1.34$ km across at $z = 650$ AU; a meter-class telescope acts as a single-pixel detector that must physically traverse the image plane, measuring the integrated brightness of the planetary Einstein ring at each sampling location and reconstructing the source by deconvolution (Turyshev & Toth 2020; Toth & Turyshev 2021; Turyshev & Toth 2022a).

The optical side of this problem is now well characterized. The aperture-averaged point-spread function falls off as $d/(4\rho)$, so each sample mixes light from the entire planet; removing this blur incurs a deconvolution penalty $\text{SNR}_R/\text{SNR}_C \simeq 0.891 D/(d\sqrt{N})$ that is strongly mitigated by wide pixel spacing (Toth & Turyshev 2021; Turyshev & Toth 2022a). The solar corona is the dominant photon background; for an exo-Earth at $z_0 = 30$ pc observed from $z = 650$ AU with a $d = 1$ m telescope at $\lambda = 1 \mu\text{m}$, the signal and corona rates are $Q_{\text{exo}} = 8.01 \times 10^4 \text{ s}^{-1}$ and $Q_{\text{cor}} = 6.20 \times 10^9 \text{ s}^{-1}$, giving a convolved-image signal-to-noise ratio $\text{SNR}_C = 43.2$ per 1800 s dwell (Turyshev & Toth 2022a; Turyshev 2026b). Megapixel-class static imaging within mission-compatible integration times follows (Toth & Turyshev 2021; Turyshev & Toth 2022a), and independent analyses of ring-based approaches (Madurowicz & Macintosh 2022) sharpened the role of coronal statistics in these budgets.

The unsolved problem is temporal. The planet rotates and orbits, its illumination changes, and—most seriously—its cloud cover reorganizes on day timescales while the raster is acquired over weeks to months. Toth & Turyshev (2023) performed the first dynamic reconstructions of a rotating, orbiting planet and recovered surface albedo only at $\text{SNR} \lesssim 2$. Toth (2025) added a stochastic cloud model, treated the clouds as noise to be averaged out (a “long-exposure” philosophy), and found that even 7–14% cloud cover degrades the recovered SNR below unity, concluding with “a cautious assessment of the practical feasibility” of SGL imaging given that Earth-like planets have $\sim 70\%$ cloud cover. The comprehensive benchmark of Turyshev (2026b) reaches a complementary conclusion: a static inverse applied to rotating cloudy data is the wrong estimator; phase-registered robust coadds recover persistent surface structure but at the cost of $\gtrsim 15$ registered visits per rotational phase bin, and a “full spin-resolved dynamic retrieval pipeline” is explicitly deferred as a next-stage deliverable. The covariance-aware biosignature framework of Turyshev (2026e) likewise argues that SGL inference must propagate structured error terms rather than treat all residuals as white. In short: the field agrees that the time-domain inverse problem is the gate through which SGL imaging must pass, and no published pipeline yet passes it under Earth-like cloud cover.

This paper attacks that gate directly, with five specific contributions. First (Sections 2–4), we build a transparent, end-to-end simulator—partially validated against the published photometric and deconvolution benchmarks—in which planetary rotation, orbital illumination, and stochastic advecting clouds

are modeled explicitly. To that first item we add the two audits the v2 framework lacked: a convergence study of the intra-dwell exposure integral (Section 4.3), which shows that the 3-node rule previously used aliases rotational harmonic $m = 36$ at unit gain, and a numerical-precision and conditioning audit of the normal equations (Sections 4.4, 4.5), which identifies the constant-albedo mode as the quantity the slot nuisance actually regularizes. Second (Section 5), we formulate reconstruction as regularized generalized least squares on time-tagged samples (time-domain inversion, TDI; distinct from hardware time-delay integration in detector readouts), with the per-dwell-slot cloud mean handled by exact anchored nuisance profiling rather than the approximate deflation of v2. We isolate the observing-cadence law under two named, mutually exclusive resource ledgers (fixed photons per raster position, or fixed wall clock) rather than the single undeclared budget of v2, and ablate the estimator — slot handling, whitening, anchoring — holding data and operator fixed. Every number is reported as a mean and standard error over paired seeds, with paired t -statistics, win counts and the exact attainable significance floor, and with a physical cloud-only null in place of v2’s geometry-scrambled surrogate. Third, we broaden the robustness tests to non-stationary latitudinal cloud banding, multi-timescale weather, per-pass independent weather, surface-coupled clouds, and a full 5×5 grid of affine cloud-climatology errors, and we state explicitly which metrics are blind to which errors (r is exactly affine invariant; bias, RMSE and contrast are not). Fourth, we extend the geometric requirement to spin-pole tilt and initial phase with measured tolerances, and we test — and refute — the v2 claim that the residual spin state is refinable in flight from the profile likelihood: the whitened objective is exactly invariant to assumed initial phase. Fifth (Section 6), we evaluate effective spatial resolution $r(\ell)$ with a real spherical-harmonic transform and an explicit quadrature error bound, land–ocean detection against two nulls with stated ensemble floors, and translate cadence requirements into mission terms using an effective-independent-look count that depends on the assumed weather decorrelation time; the target table is a reproducible nominal-proxy list, not an uncertainty propagation.

2 SGL Photometric Model and Validation

2.1 Optical response

We adopt the standard monopole, aperture-averaged description of SGL image formation (Turyshev & Toth 2020; Toth & Turyshev 2021; Turyshev 2026b). A source-plane coordinate ξ maps to the image-plane coordinate $\rho \simeq -(z/z_0) \xi$; an Earth-radius planet at $z_0 = 30$ pc observed from $z = 650$ AU projects

to

$$D_{\text{img}} = 2R_p \frac{z}{z_0} = 1.338 \text{ km.} \quad (1)$$

A telescope of aperture d used as a ring photometer averages the J_0^2 oscillations of the monopole PSF; the effective kernel is

$$K(0) = 1, \quad K(\rho > 0) \simeq \frac{d}{4\rho}, \quad (2)$$

renormalized so that the mean convolved signal over the projected disk of a reference (albedo 0.3, fully illuminated) planet is unity (Turyshev 2026b). Our discrete kernel follows the $d/4\rho$ tail over the full grid (Fig. 1a). Because K falls only as $1/\rho$, every sample mixes light from the entire disk—the property that couples the temporal variability of any part of the planet into every measurement.

2.2 Photon rates and noise

We normalize the scene such that a value of 1 corresponds to the disk-mean signal of the reference planet, which maps to the fiducial rates of Turyshev & Toth (2022a): $Q_{\text{exo}} = 8.01 \times 10^4 (d/1 \text{ m})^2 (650 \text{ AU}/z)^{1/2} (30 \text{ pc}/z_0) (\lambda/1 \mu\text{m}) \text{ s}^{-1}$ and $Q_{\text{cor}} = 6.20 \times 10^9 (d/1 \text{ m})^2 (650 \text{ AU}/z)^2 (\lambda/1 \mu\text{m}) \text{ s}^{-1}$ after coronagraphic rejection and annular extraction, with the mean corona assumed subtracted and only its shot noise retained. Each dwell of duration t_s at image-plane position ρ_k yields a Gaussian sample with standard deviation

$$\sigma_k = \frac{\sqrt{(Q_{\text{cor}} + Q_{\text{exo}} \bar{s}_k) t_s}}{Q_{\text{exo}} t_s}, \quad (3)$$

where $\bar{s}_k$ is the convolved scene value. For $t_s = 1800 \text{ s}$ this gives $\text{SNR}_C = 1/\sigma = 43.16$, matching Turyshev (2026b) exactly within the adopted photon-noise assumptions.

2.3 Deconvolution-penalty validation and numerical limits

As a check of the discrete operator we reproduce the deconvolution noise-amplification analysis of Toth & Turyshev (2021). For a static, fully illuminated planet we convolve, add white noise, deconvolve by Fourier quotient, and measure the ratio $\text{SNR}_R/\text{SNR}_C$ of reconstructed noise to input noise. Figure 2 compares measurements with the analytic formula $0.891 D/(d\sqrt{N})$. For $n = 32$ ($D = 41.8 \text{ m}$) we measure 0.657 vs. 1.0 (saturating toward unity); for $n = 64$ ($D = 20.9 \text{ m}$), 0.324 vs. 0.291; for $n = 128$ ($D = 10.5 \text{ m}$), 0.106 vs. 0.073.

The $\sim 45\%$ difference at $n = 128$ has a clear physical and numerical origin: the continuous formula $0.891 D/(d\sqrt{N})$ assumes continuous circular integration of the $1/\rho$ Fourier transform and high-frequency asymptotes on an infinite domain. On a discrete numerical grid with finite pitch D , pixel-binning across the central $1/\rho$ singularity acts as an effective low-pass filter, and zero-padded circular boundary conditions attenuate high-wavenumber noise amplification relative to the continuum integral. We emphasize this constitutes partial validation of the

discrete aperture-averaged forward operator rather than full optical validation of an in-flight SGL mission.

3 Dynamic Scene and Campaign Model

3.1 Rotating, illuminated planet

The truth surface is a seeded synthetic albedo map (ocean 0.06, continents 0.32 ± 0.10 , polar caps up to 0.60; 30% land fraction) generated by thresholded multi-octave Gaussian random fields on a 144×72 longitude–latitude grid (Fig. 1b). The planet rotates with period $P_{\text{rot}} = 24 \text{ h}$ and is illuminated by its host star with a Lambertian terminator whose sub-stellar point drifts through 86.1° of orbital phase over the **87.33-day fiducial campaign** ($P_{\text{orb}} = 1 \text{ yr}$, $360^\circ \times 87.33/365.25$; 88.7° over a **90.0-day window**). The simulation holds the sub-stellar latitude fixed at the equator: obliquity is not modeled, so the illumination drift here is purely orbital in longitude and the seasonally-migrating intertropical zone of a tilted planet is outside the tested regime (Section 7). For the short-period habitable-zone planets around M dwarfs listed in Section 6.1 ($P \sim 10\text{--}20 \text{ d}$), a 90-day campaign spans 5–9 full orbits and therefore samples all orbital phases repeatedly; the single sub- 90° drift modeled here is the *least phase-covering case*, not a conservative envelope for every target, because the two regimes differ in which longitudes are ever illuminated (a synchronous, zero-obliquity planet keeps a fixed sub-stellar longitude). We therefore report the tested case and do not claim the benchmark is conservative for all listed targets. Each dwell is integrated over $n_{\text{sub}} = 16$ Gauss–Legendre sub-exposures spanning the dwell, so intra-dwell rotational smear (7.5° per 1800 s dwell) is present in both the synthetic data and the forward operator. A uniformly-spaced 3-node rule was tested and rejected: its node spacing of $t_s/3$ aliases azimuthal harmonic $m = 36$ onto the disk-integrated mean with unit gain instead of annihilating it, a response error of order unity that propagates directly into the reconstruction (Section 4.3).

3.2 Stochastic cloud fields: stationary and structured

The baseline cloud model is an advecting Ornstein–Uhlenbeck (OU) Gaussian random field on the fine grid: spatial correlation length 12° , zonal advection 6° day^{-1} , temporal decorrelation $\tau_c = 4 \text{ d}$, mapped through a soft threshold to an opacity screen $c \in [0, 1]$. The effective scene is $A_{\text{eff}} = (1 - c)A_{\text{surf}} + cA_{\text{cl}}$ with $A_{\text{cl}} = 0.65$. **The sweep is over the specified cover parameter $f_c \in \{0, 0.25, 0.40, 0.55, 0.70\}$; because the soft threshold saturates, the realized mean covers are slightly different (0.22, 0.38, 0.55, 0.70 respectively, averaged over seeds) and every table in this paper quotes specified f_c together with the realized cover it produced.**




Figure 1: Model overview. (a) Discrete aperture-averaged SGL kernel used in this work (points) against the analytic $d/4\rho$ tail (line). (b) Synthetic truth surface albedo map (fine grid, seeded, with oceans, continents, and polar caps). (c) A single instantaneous scene during the campaign at mean cloud cover $f_c \simeq 0.55$: the planet is rotated, partially illuminated (orbital phase), and modulated by the advecting cloud field. (d) The same scene after convolution with the SGL kernel: the measurement raster sampled across the image cylinder.

To test robustness against non-terrestrial weather we implement four extended climatology variants, each re-normalized so that the realized mean cover matches the baseline f_c it is compared with (otherwise a variant that simply has more or less cloud would be confounded with a variant that has a different cloud *pattern*): (i) *Banded zonal profile*: the opacity field is multiplied by $w_{\text{lat}}(\theta) = 1 + 0.35 \cos 2\theta - 0.25 \cos 4\theta$, normalized to unit mean. We previously labelled this an “ITCZ” profile, which it is not: its maximum (1.31) lies at $\pm 35^\circ$ — the subtropical region an ITCZ model should *suppress* — its equatorial value is a mild 1.10, and it is symmetric about the equator and never migrates. It is a test of latitudinal inhomogeneity only, and we report it as such. (ii) *Surface-correlated cloudiness*: coupling cloud opacity to surface reservoirs, $g_{\text{eff}} = g - \beta_{\text{surf}}(A_{\text{surf}} - \bar{A})$, enhancing convection over low-albedo tropical oceans; and (iii) *Multi-timescale temporal evolution*: summing a rapid convective mode and a persistent synoptic regime with weights (0.3, 0.7) and correlation times (1 d, 10 d), nor-

malized to unit variance. (iv) *Per-pass independent weather*: a fresh, temporally uncorrelated cloud realization for every pass, included as the causal control of Section 5.5 rather than as a physical climate.

3.3 Campaign and sampling geometry

The campaign matches current swarm concepts (Turyshev & Toth 2022a; Helvajian et al. 2023): $N_{\text{sc}} = 16$ spacecraft sample an $n \times n = 64 \times 64$ raster ($\Delta_{\text{img}} = 20.9$ m pitch, 4096 pixels) in boustrophedon order, divided into 256 dwell slots. In each dwell slot, 16 pixels are sampled simultaneously by spacecraft spaced across adjacent raster tracks with transverse separations of ~ 80 –300 m. **Nominal dwells are $t_s = 1800$ s with $M_p = 16$ complete passes, giving 65,536 time-tagged samples, 8.00 h of integration per raster position, 85.33 d of exposure time, 2.13 d of charging overheads, and a wall-clock length of 87.33 d — not the 90.0 d quoted previously, which silently assumed zero overhead. Inter-pass dead time is now an ex-**




Figure 2: Validation of the discrete operator: measured deconvolution noise penalty vs. the analytic estimate of Toth & Turyshv (2021). The $\sim 45\%$ difference at $n = 128$ reflects discrete pixel-averaging of the central singularity and finite-domain boundary attenuation.

PLICIT, AUDITABLE LEDGER INSTEAD OF A RANDOM DRAW: the fiducial campaign charges $T_{\text{oh}} = 45$ s per dwell (30 s translation slew, 10 s pointing settling, 5 s inter-spacecraft metrology) and no additional random gap. The earlier 0–12 h random gaps are removed because they made the wall-clock length seed-dependent and were never reconciled with the 90-day claim.

3.4 Resource conventions: two mutually exclusive ledgers

The referee correctly observed that “fixed 8 h per pixel” and “fixed 90-day campaign” cannot both be satisfied as M_p grows. We therefore run the cadence study under two explicitly named arms and never mix them.

- **Arm A (photons):** hold $M_p t_s = 8$ h per raster position and let the wall clock be whatever it needs. Then $t_s = 7200, 3600, 1800, 900, 450$ s for $M_p = 4, 8, 16, 32, 64$ give 85.73, 86.27, 87.33, 89.47, **93.73 d**; the $M_p = 64$ configuration *violates* a 90-day cap.
- **Arm B (wallclock):** force the campaign to 89.87 d including overheads and let t_s fall to 7548.8, 3751.9, 1853.4, 904.2, **429.6** s, so the photons per position fall from 8.39 h to 7.64 h instead of staying at 8 h.

Every cadence number in Section 5.5 is reported per arm, and the two arms answer different questions: Arm A isolates the statistical effect of revisit count at constant photons, Arm B is what a 90-day mission can actually buy.

Cloud-induced signal fluctuations are calibrated against the data the estimator actually sees, i.e. the per-sample standard deviation of the cloud-only component of the disk-integrated signal with the 16-node Gauss–Legendre exposure integration of Section 4.3. The resulting values are $\sigma_{\text{cl}} = 0.296, 0.378, 0.441, 0.486$ for $f_c = 0.25, 0.40, 0.55, 0.70$, against the coronal single-dwell noise $\sigma_{\text{cor}} = 0.0232$ at $t_s = 1800$ s. That is a standard-deviation ratio of 12.8–21.0, or a *variance* ratio of 163–440. The previously quoted “14–23 times” was a standard-deviation ratio computed with the rejected 3-node quadrature; it is neither the variance ratio its wording implied nor reproducible under the corrected operator. Planetary weather dominates the single-sample photometric error budget by more than two orders of magnitude in variance.

4 Reconstruction Methods

4.1 Time-domain inversion (TDI)

Two notational points precede the estimator. We write the target distance as z_0 (pc-scale) and the spacecraft heliocentric distance as z (AU-scale), reserving z/z_0 for the imaging magnification; the minimum heliocentric distance for a ray of impact parameter b is $z_{\text{min}} \simeq b^2/2r_g$ and is *not* the same quantity as the target distance z_0 . Second, all reconstructions use $n_{\text{sub}} = 16$ Gauss–Legendre sub-exposures (Eq. 4); the number of nodes is a parameter of the operator, not a free-tuning knob, and Section 4.3 shows why the previously used 3-node rule is inadequate. Let $\mathbf{s} \in \mathbb{R}^{n_s}$ be the static surface albedo map on a 72×36 reconstruction grid ($n_s = 2592$, coarser than the truth grid). Each sample k , taken at time t_k and image-plane pixel p_k , satisfies

$$y_k = \mathbf{f}_k^T \mathbf{s} + \varepsilon_k, \quad \mathbf{f}_k^T = \frac{1}{n_{\text{sub}}} \sum_{m=1}^{n_{\text{sub}}} \mathbf{k}_{p_k}^T \Pi(t_k + \delta t_m), \quad (4)$$

where $\Pi(t)$ renders the map at the rotation phase, illumination, and spin-pole geometry of time t , and $\mathbf{k}_{p_k}$ is the SGL kernel row. Setting $n_{\text{sub}} = 1$ corresponds to an instantaneous midpoint model; $n_{\text{sub}} = 16$ uses the **Gauss–Legendre nodes δt_m and weights of the dwell interval, which is what the synthetic data are integrated with.** The error ε_k contains coronal shot noise and cloud fluctuations. The climatological mean-cloud bias is removed by the affine correction $\hat{\mathbf{s}} = (\hat{\mathbf{s}}_{\text{raw}} - f_c A_{\text{cl}})/(1 - f_c)$.

TDI solves the regularized generalized least-squares problem:

$$\hat{\mathbf{s}} = \arg \min_{\mathbf{s}} \|\mathbf{C}^{-1/2}(\mathbf{y} - \mathbf{F}\mathbf{s})\|^2 + \lambda_{\text{eff}} \|\mathbf{L}\mathbf{s}\|^2, \quad (5)$$

with spherical-grid Laplacian $\mathbf{L}$ and $\lambda_{\text{eff}} = \lambda \text{tr}(\mathbf{F}^T \mathbf{C}^{-1} \mathbf{F}) / \text{tr}(\mathbf{L}^T \mathbf{L})$ ($\lambda = 3 \times 10^{-3}$, frozen on calibration seed 99). That weight is carried unchanged through the revised suite; it is not claimed to be optimal — Section 5.13 sweeps six decades of it and finds the objective unable to select one without truth. $\mathbf{C}$ is the working covariance, block

diagonal over image-plane pixels: the M_p samples of pixel p (one per pass, at the same raster position) are correlated by the OU cloud field,

$$(C_p)_{ij} = \sigma_{\text{cl}}^2 \exp(-|t_i - t_j|/\tau_c) + \delta_{ij} \sigma_i^2, \quad (6)$$

while different pixels are treated as conditionally independent given the surface map (the spatial cloud correlation is absorbed into σ_{cl} rather than modeled across pixels; Section 7 states what this approximation costs). Four weighting/projection combinations are ablated in Section 5.4: (White) diagonal $C = \text{diag}(\sigma_k^2 + \sigma_{\text{cl}}^2)$, the assumption that clouds average like uncorrelated white noise (Toth 2025); (OU) Eq. 6 without any slot treatment; (v2 deflation), the approximate per-slot projection described below; and (Full TDI), exact nuisance profiling with Eq. 6.

4.2 Nuisance profiling of the simultaneous-sample offsets

Because $K(\rho) \propto 1/\rho$, an instantaneous disk-integrated cloud fluctuation perturbs all $N_{\text{sc}} = 16$ simultaneous samples coherently. The v2 treatment subtracted the per-slot mean of both y and F and rescaled the marginal noise by $\sqrt{1 - 1/N_{\text{sc}}}$, claiming that this “removes 256 common modes (6.25% rank reduction) while retaining $1 - 1/N_{\text{sc}} = 93.75\%$ of the information”. Both the mechanism and the numbers in that sentence are wrong, and the revised treatment replaces it with the exact profile likelihood of the referee’s eq. (R11).

Write the model with slot nuisance offsets α ,

$$y = Fs + B\alpha + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, C), \quad (7)$$

where B is the indicator matrix of dwell slots (B has one column per slot, 4096 columns). The key structural fact, which the v2 “orthogonal projection” narrative missed, is that C is block diagonal over *image-plane pixels*, while B couples slots *within* a pixel. The M_p samples belonging to pixel p sit in M_p different slots, so

$$H \equiv B^T C^{-1} B = \text{blkdiag}(H_1, \dots, H_{n_{\text{bin}}}), \quad (8)$$

with $n_{\text{bin}} = 256$ blocks of size $M_p \times M_p$ — *not* diagonal, i.e. the offsets are statistically coupled through the shared temporal cloud covariance of a single pixel. Profiling them out therefore requires inverting those $M_p \times M_p$ blocks: with $G_j = B_j^T C^{-1} F$, $\mathbf{g}_j = B_j^T C^{-1} y$ and $H_j = B_j^T C^{-1} B_j$, the free (unconstrained) profile is

$$A_{\text{free}} = F^T C^{-1} F - \sum_j G_j^T H_j^{-1} G_j, \quad (9)$$

$$\mathbf{b}_{\text{free}} = F^T C^{-1} y - \sum_j G_j^T H_j^{-1} \mathbf{g}_j. \quad (10)$$

With the standard anchoring constraint $\sum_b \alpha_b = 0$ this collapses to a rank-one correction,

$$A = A_{\text{free}} + \mathbf{q}\mathbf{q}^T/S, \quad \mathbf{b} = \mathbf{b}_{\text{free}} + \mathbf{q}m/S, \quad (11)$$

where $\mathbf{v}_j = H_j^{-1} \mathbf{1}$ and

$$\mathbf{q} = \sum_j G_j^T \mathbf{v}_j, \quad S = \sum_j \mathbf{1}^T \mathbf{v}_j, \quad m = \sum_j \mathbf{v}_j^T \mathbf{g}_j. \quad (12)$$

which is exact and costs only the block inversions already required by Eqs. 9–10. Without anchoring the offsets are only projected out, and the estimator loses the disk-integrated albedo information entirely; Eq. 11 is what “full TDI” means in this paper.

4.3 Exposure quadrature: convergence of the smear operator

Eq. 4 averages the illumination/rotation operator over the dwell, and the synthetic data are generated with the same rule, so the sub-exposure count n_{sub} is part of the operator definition rather than a tuning knob. v2 used three *equally spaced*, equal-weight sub-exposures at $-t_s/3$, 0 , $+t_s/3$, which is neither a Gauss–Legendre rule nor a convergent quadrature of the smear kernel. Writing the exact rotational-harmonic response of a dwell of length t_s as $R_m(t_s) = \text{sinc}(m\Omega t_s/2)$ with $\Omega = 2\pi/P_{\text{rot}}$, the equal-spaced rule gives

$$R_m^{(\text{eq3})} = \frac{1}{3} \left[2 \cos\left(\frac{m\Omega t_s}{3}\right) + 1 \right], \quad (13)$$

so at $t_s = 7200$ s the exact response $R_{36} = 0$ becomes $R_{36}^{(\text{eq3})} = 1.000000$. At $P_{\text{rot}} = 24$ h the node spacing $t_s/3$ is 2400 s at $t_s = 7200$ s, i.e. 10° of longitude, which is exactly the period of the $m = 36$ azimuthal harmonic on the 72-column grid. That harmonic is sampled at its own period, so instead of being averaged to zero it is passed at unit gain; the same holds approximately for $m = 35$ –40, where the equal-spaced rule returns $+0.99$ to $+0.84$ against true responses of ± 0.07 . The worst error over $m \leq 40$ at $t_s = 7200$ s is 1.02 — larger than the response itself. This is a quadrature alias, not a smoothing error: the rule is blind to the very structure it is meant to average out, and data and operator generated with it are self-consistent, so the mismatch is invisible in residuals and shows up only in the recovered map.

Table 1 reports $\max_{m \leq 40} |R_m^{(n_{\text{sub}})} - R_m|$ for Gauss–Legendre rules, computed by `sglsim.exposure_quadrature_convergence`. Convergence is monotonic in n_{sub} at fixed t_s and the GL rules never recover unit gain at an aliased harmonic: GL-3 peaks at 0.88 rather than 1.02, and GL-16 is *spectrally converged* — its worst error over all five dwell lengths is 2.6×10^{-13} , i.e. machine precision even for the 7200 s dwells of Arm A, where the intra-dwell smear alone is 30° . GL-16 is therefore used everywhere: in the data, in the forward operator, in the σ_{cl} calibration of Section 3.4, and in every reconstruction. The practical consequence is that the $t_s = 7200$ s ($M_p = 4$) cadence points of the original analysis were computed with an operator whose response error at the highest retained harmonic is of order unity; they are superseded by the Arm A/Arm B numbers of Section 5.5.

Table 1: Maximum rotational-harmonic response error over $1 \leq m \leq 40$ against the exact sinc, by dwell length t_s . The top row is the equal-spaced rule used in v2; the rest are Gauss–Legendre. From results/quadrature.npz (eq3_max_err, gl_max_err).

Rule	$\max_{m \leq 40} R_m^{(n)} - R_m $ by dwell t_s				
	7200	3600	1800	900	450 s
equal-3 (v2)	1.0	1.3×10^{-1}	3.6×10^{-2}	2.4×10^{-2}	7.4×10^{-3}
GL-3	8.8×10^{-1}	2.7×10^{-1}	8.3×10^{-3}	1.5×10^{-4}	2.5×10^{-6}
GL-5	5.1×10^{-1}	3.5×10^{-3}	5.3×10^{-6}	5.8×10^{-9}	5.8×10^{-12}
GL-8	4.8×10^{-3}	2.4×10^{-7}	4.9×10^{-12}	4.4×10^{-16}	2.2×10^{-16}
GL-16	2.6×10^{-13}	6.4×10^{-16}	6.4×10^{-16}	2.2×10^{-16}	2.2×10^{-16}
GL-32	1.2×10^{-15}	1.1×10^{-15}	1.0×10^{-15}	4.4×10^{-16}	2.2×10^{-16}

4.4 What profiling actually removes: an information audit

The rank/percentage framing of the v2 deflation paragraph cannot be repaired, so we replaced it with a direct eigenvalue audit of the normal matrix at the fiducial campaign (Table 2, produced by `sglsim.profiling.information.audit`). Three findings.

First, profiling removes no rank at all. The un-profiled normal matrix already has nullity 288, the free profile has nullity 288, and the anchored profile has nullity 288. The null space is a property of the reconstruction grid and the raster, not of the nuisance treatment: 144 basis columns are *identically zero* because they live in the two polar latitude rows of the equirectangular grid, which a circular image-plane raster never samples (the projected disk does not reach $|\text{lat}| = 90^\circ$ cell centres). The remaining 144 null directions are the associated near-degenerate partners; the smallest retained eigenvalue with λ_{eff} regularization is 6.3×10^{-6} against a maximum of 2.9×10^3 .

Second, what profiling removes is the *response along the mean-albedo direction*, measured as the constant-mode gain $\|A1\|/\|1\|$: 913.7 un-profiled, 18.79 for the free profile (2.06% retained), and 774.1 for the anchored profile (84.72% retained). The free profile is therefore not a mild correction — it attenuates the disk-integrated albedo by a factor of 49, which is why anchoring is mandatory rather than cosmetic.

Third, the identity $W_\perp B = 0$ quoted in v2 is a trivial property of any profile likelihood (the residual is orthogonal to the nuisance space by construction); it is not evidence of a surface-mode degeneracy, and an earlier version of this work read it as one.

4.5 Numerical precision verification

The normal matrix $A = F^T C^{-1} F$ is accumulated in single precision (‘float32’) and solved in double precision (‘float64’). The v2 precision test was vacuous: it compared two double-precision accumulations and reported their difference, which measures nothing. It is replaced by a test that actually probes the single-precision path — the same design matrix, rounded to float32 before accumulation, ver-

Table 2: Information audit of slot-offset profiling at the fiducial campaign ($n = 64$, $N_{\text{sc}} = 16$, $M_p = 16$, $n_s = 2592$). Nullity is unchanged by profiling; the constant-mode gain is what changes. free = offsets projected out, anchored = Eq. 11.

Quantity	un-profiled	free	anchored
nullity	288	288	288
exactly-zero basis columns	144	144	144
constant-mode gain	913.7	18.79	774.1
fraction of gain retained	1	0.0206	0.8472
largest eigenvalue	2940	429.9	1030
λ_{eff}	0.00842	0.00558	0.00564

sus the float64 reference — at a reduced but honest problem size ($n = 32$, 16,384 samples; the fiducial $n = 64$ design would need 2.8 GB in two float64 copies inside the test harness), with the results in Table 3. Three numbers matter. The float32 design matrix differs from the float64 reference by 3.1×10^{-7} in Frobenius norm and the right-hand side by 4.1×10^{-7} , i.e. at the float32 machine-precision level, and the same quantities differ by only 1.0×10^{-8} and 1.5×10^{-8} from rounding alone — confirming the discrepancy is representational, not algorithmic. Second, changing the *summation order* of the blockwise float64 accumulation perturbs A by 5.3×10^{-16} , which is the number the v2 test was reporting. Third, and most relevant for the science: solving the regularized normal equations versus a QR solve of the same whitened design changes the recovered map by 1.7×10^{-13} relative to the map norm, and chunking the profiler in different orders changes the fiducial map by 3.4×10^{-13} . The unregularized normal matrix is exactly singular (rank 2016 of 2592 at the reduced size, nullity 288 at the fiducial size, Table 2), so $\kappa(A)$ is reported only after regularization, where $\kappa = 5.3 \times 10^5$ and the normal-equation residual is 1.3×10^{-15} . Single-precision accumulation is therefore adequate for this estimator, but the claim is now stated with its scope: it is a statement about the linear algebra, verified at reduced size, not a general bound.

Table 3: Numerical precision audit (sglsim.check_numerical_precision, $n = 32, N_{\text{sc}} = 16, M_p = 16$).

Quantity measured	Relative error
float32 vs float64 design A (Frobenius)	3.1×10^{-7}
float32 vs float64 right-hand side b	4.1×10^{-7}
rounding-only contribution to A	1.0×10^{-8}
float64 summation-order change to A	5.3×10^{-16}
normal-equation vs QR solve (map)	1.7×10^{-13}
profiler chunk-order change (map)	3.4×10^{-13}
regularized condition number $\kappa(A)$	5.3×10^5
unregularized rank / grid size	2016/2592

4.6 Baselines and statistical metrics

We compare TDI against two registration-free coadd baselines, both evaluated on the same surface grid and with the same debiasing. B1 is the phase-blind limit of the same pipeline (a single rotational bin: the all-visit mean raster, Wiener-deconvolved and back-projected through the mean illumination; cf. Toth & Turyshev 2021). B2 is the phase-registered coadd baseline emulating Turyshev (2026b): samples are binned in rotational phase, each bin is Wiener-deconvolved, and the bins are back-projected onto the sphere by bilinear interpolation weighted by illumination.

The bin count of B2 is not a free parameter and we now fix it explicitly. A pass visits a given image-plane pixel once, and consecutive dwell slots are $t_s = 1800$ s apart, i.e. $1/48$ of the rotation period: there are 48 distinct sample phases per rotation. Coadding into only 8 phase bins averages over 45° of longitude, and on ideal cloud-free noiseless data that baseline collapses to $r = 0.07$ for reasons that have nothing to do with time-domain inversion; $r = 0.94$ is obtained at 16 bins, which is the best-tuned compromise between rotational smear and the number of samples per bin ($3 \times M_p = 16$ here). All B1/B2 numbers quoted below use $n_{\text{bins}} = 16$ for B2 and the same Wiener regularization $K_w = 3 \times 10^{-3}$; the naive 8-bin variant is reported in Table 4 only to quantify the artifact, so that the competitor is not a straw man.

Metrics are evaluated over latitudes $|b| \leq 60^\circ$ across an expanded ensemble of 10 paired seeds (seeds 11–20). Scalar metrics (r , SSIM, NRMSE) are computed per realization seed and then averaged over seeds (with 95% bootstrap confidence intervals), rather than computed on a seed-averaged map. Because 10 seeds is a small ensemble, we now report the exact attainable significance floor alongside every p -value: for $n = 10$ paired samples the smallest attainable two-sided sign-flip (Wilcoxon) p is $2/2^{10} = 1.95 \times 10^{-3}$, so no paired test on this ensemble can support $p < 10^{-3}$, and any such statement in v2 was not attainable. We also report a second, independent null: a cloud-only null in which the surface is replaced by the mean and only the weather is retained, whose finite ensemble size limits attainable p to

$1/(N_{\text{null}} + 1)$. Finally, SSIM and the land–ocean contrast ratio are computed on the coarse 36×72 reconstruction grid with the same latitude mask as r , and the harmonic decomposition of Section 5.2 uses real spherical harmonics with $\cos \vartheta$ quadrature weights, not a 2D FFT on an equirectangular array (which is not a spherical-harmonic decomposition: the grid cells are not equal-area and the azimuthal sampling is aliased at high $|\text{lat}|$).

5 Results

5.1 Planetary rotation under characterized geometry

Headline result, cloud-free. For a rotating, orbit-illuminated cloud-free planet observed with the fiducial 90-day campaign and exactly known spin geometry, the profiled time-domain estimator recovers Pearson $r = 0.9786 \pm 0.0002$ (standard error over 10 paired weather seeds), SSIM = 0.8482 ± 0.0010 , normalized RMSE 0.0628 ± 0.0003 , albedo bias -0.0006 ± 0.0006 , and land–ocean contrast ratio 1.056 ± 0.001 (Fig. 3, top center; Table 4, first column). All numbers in this section and in Table 4 come from the 16-node exposure quadrature of Section 4.3, the exact anchored nuisance profiling of Section 4.2, and phase bins registered into 16 rotational bins; they therefore differ from the v2 values ($r = 0.990$, SSIM = 0.918), which used the approximate deflation and a coarser quadrature.

To attribute what remains, we compare three reference configurations on the *same* data. (i) The phase-blind coadd (B1), which treats the campaign as a single static exposure, collapses to $r = 0.1155 \pm 0.0004$: with rotation unmodeled the surface is essentially unrecoverable, and the $+0.863$ gap to TDI is the contribution of the time-registered forward operator alone. (ii) Registering the coadds into 16 rotational phase bins (B2) recovers most of the signal, $r = 0.9361 \pm 0.0001$, but with a -0.048 albedo bias and a contrast ratio of only 0.844: binning averages within each 22.5° phase window and flattens the map. (iii) A white-noise GLS fit with the same time-dependent operator reaches $r = 0.9802 \pm 0.0002$, i.e. 0.0015 above the profiled estimator (paired difference -0.0015 , 0/10 seeds in favour of whitening, $t = -22.6$).

Finally, the static ideal-data reference the referee asked for. We rebuild the campaign with the planet *frozen*: the same 64×64 raster, the same 16 revisits and the same 1800 s dwell — so every raster position receives exactly the same 8 h of integration and there is no weather — but all passes see one identical pose, so there is no rotational smearing to model and no phase coverage to solve. Scored on the same seeds with the same estimator (archive static_reference.npz), that best case of any observation of this planet at this dose gives $r = 0.4074 \pm 0.0254$ (standard deviation across seeds 0.0804, per-seed range 0.284–0.523), SSIM = 0.3604, normalized RMSE 0.4299. It is

worse than the rotating cloud-free planet by $+0.5728 \pm 0.0254$ in favour of rotation ($p = 3 \times 10^{-9}$, 10/10 seeds), and the rotating cloudy campaign sits between them: -0.6379 ± 0.0083 relative to rotating cloud-free ($p = 5 \times 10^{-14}$, 10/10). The reason is the conditioning of the operator rather than any information argument: with the planet frozen, the 16 pose-passes of each raster position are literally identical rows of F , so the time-domain design loses the rank that rotation supplies and the regularizer does all the work — a frozen planet is not an easier inverse problem, it is a degenerate one. This is the quantitative form of the claim above: the loss attributable to rotation *per se* is not a loss at all, and the whole 0.638 gap the paper is about is weather, measured against a rotation reference that rotation is helping.

That last comparison is the honest statement of what rotation alone costs. With no weather there is no correlated noise for temporal whitening to remove, so whitening buys nothing and slightly dilutes the fit; the remaining signal correlation comes from the time-dependent operator, which both configurations share. Rotation *per se* is therefore not the information bottleneck: the scale-dependent correlation $r(\ell)$ of Section 5.3 stays above 0.99 for every degree $\ell \leq 20$ the grid supports, so no rotational band is lost relative to the $\ell = 20$ quadrature ceiling, and the residual $1 - r \simeq 0.021$ is regularization and aperture, not aliasing. The v2 claim that rotation “introduces virtually zero information loss” was stronger than these data support: rotation is *tractable* under characterized geometry, and the degradation budget is dominated by weather, not spin.

5.2 Cloud degradation and effective spatial resolution

Table 4 and Figure 4 document the collapse of fidelity with cloud cover at fixed single-campaign resources. The cover realized by the OU field differs from the nominal parameter, so both are quoted: mean realized cover is 0.237, 0.398, 0.558 and 0.710 for the nominal $f_c = 0.25, 0.40, 0.55, 0.70$ runs.

Mean correlation falls from 0.9786 (cloud-free) to 0.5966 ± 0.0083 at 23.7% realized cover, 0.4538 ± 0.0070 at 39.8%, 0.3422 ± 0.0078 at 55.8% and 0.2491 ± 0.0100 at 71.0%. SSIM falls faster: 0.288 ± 0.005 already at 23.7% cover and 0.061 ± 0.003 at 71.0%, i.e. small-scale structure is erased first, which is the expected signature of a noise process whose power is broad in ℓ (Section 5.3).

Two comparisons in Table 4 qualify the headline claim. (1) At every cloudy cover the profiled estimator beats the white-noise GLS fit by $\Delta r = +0.069$ to $+0.054$ (10/10 paired seeds in favour, $t = 9.3$ to 7.4), so the OU temporal covariance model earns its keep only once weather is present. (2) At 23.7% cover the 16-bin phase-binned coadd (B2) actually scores *higher* than TDI, $r = 0.6097$ vs. 0.5966 (paired $\Delta r = -0.013$, 1/10 seeds in favour of TDI, $t = -2.62$), and the gap only reverses at higher cover ($+0.006$ at 39.8%,

Table 4: Map fidelity vs. cloud cover for the fiducial campaign ($n = 64$, $M_p = 16$, $t_s = 1800$ s, 16-node quadrature, exact anchored profiling). Rows are means $\pm$ standard error over 10 paired weather seeds; the B2/B1 rows use 16, 8 and 1 rotational phase bins respectively.

Method	$f_c=0$	0.24	0.40	0.56	0.71
Pearson r					
TDI (profiled)	0.9786 ± 0.002	0.5966 ± 0.0083	0.4538 ± 0.0070	0.3422 ± 0.0078	0.2491 ± 0.0100
White GLS (B3)	0.9802 ± 0.002	0.5274 ± 0.0080	0.3863 ± 0.0076	0.2886 ± 0.0092	0.2107 ± 0.0102
Phase-binned (B2)	0.9361 ± 0.001	0.6097 ± 0.0079	0.4479 ± 0.0115	0.3175 ± 0.0150	0.2028 ± 0.0149
Phase-blind (B1)	0.1155	0.0782	0.0567	0.0417	0.0328
B2, 8 bins	0.0700	0.0136	-0.0140	-0.0309	-0.0416
SSIM					
TDI (profiled)	0.8482	0.2876	0.1748	0.1060	0.0606
White GLS (B3)	0.8568	0.2492	0.1447	0.0853	0.0472
Phase-binned (B2)	0.7947	0.3515	0.2223	0.1307	0.0627
Phase-blind (B1)	0.0831	0.0376	0.0143	0.0028	-0.0011
albedo bias / contrast ratio					
TDI (profiled)	-0.001/1.056	-0.017/1.103	-0.009/1.107	+0.004/1.071	+0.019/1.024
Phase-binned (B2)	-0.048/0.844	-0.057/0.862	-0.034/0.851	+0.003/0.826	+0.065/0.786

$+0.025$ at 55.8%, $+0.046$ at 71.0%, $t = 0.79$ to 6.14). B2 is a non-parametric estimator: it makes no assumption about cloud statistics, so at low cover — where the cloud nuisance is small — its binning variance costs less than TDI’s model error. The claim we can defend is therefore not “TDI is always better than phase-binning,” but “TDI’s advantage grows with cover, and at high cover B2 additionally carries a bias that TDI does not” ($+0.065$ vs. $+0.019$ albedo bias and contrast ratio 0.786 vs. 1.024 at 71% cover).

5.3 Effective spatial resolution and the land-ocean test

Scale-dependent correlation. Decomposing the maps in real spherical harmonics with $\cos \vartheta$ quadrature weights over degrees $\ell \in [1, 20]$ (Fig. 6c) gives, over 10 seeds: $r(\ell) = 1.000 \pm 0.001$ averaged over $\ell = 1-4$, 0.9995 over 5-8, 0.9988 over 9-12, 0.9942 over 13-16 and 0.9924 over 17-20 for the cloud-free campaign; the cloudy campaign (55.8% realized cover) gives 0.700, 0.594, 0.418, 0.244 and 0.263 in the same bands. The v2 statement that the cloud-free curve “extends past $\ell = 20$ ” cannot be tested on this grid: 36×72 pixels cannot represent degrees above $\ell \approx 18$, and the discrete-basis orthonormality error of the quadrature itself is 0.018 at $\ell = 20$, so we quote $\ell \leq 20$ as the measurement ceiling and report the ceiling error alongside. The cloud-free result should therefore be read as “no rotational band within the supported range is lost,” not as a resolution claim. For the cloudy case the curve is not monotone at the seed level (the per-degree scatter is 0.05–0.12), so a single crossing degree ℓ_{eff} is fragile: the raw curve dips below 0.5 at $\ell = 8$ for some seeds and $\ell \in \{1, 9\}$ for others. The defensible statement is a banded one: a five-point running mean of $r(\ell)$ crosses 0.5 at $\ell \simeq 9$ (half-wavelength $\pi R_{\oplus} / \ell \approx 2200$ km for an Earth-radius planet), only 50% of seeds have $r(\ell) \geq 0.5$




Figure 3: Reconstruction gallery for the fiducial 90-day campaign ($n = 64$, $M_p = 16$, $t_s = 1800$ s, 16-node exposure quadrature). Top: truth surface albedo; TDI reconstruction of the rotating cloud-free planet ($r = 0.979$, $\text{SSIM} = 0.848$); phase-blind coadd (B1) at realized cover 0.56, which is essentially featureless ($r = 0.042$). Bottom, all at realized cover $\simeq 0.56$: 16-bin phase-binned coadd (B2), white-noise GLS (B3), and TDI. The gross continental dichotomy is recovered, but fine features are erased by weather.

at $\ell \geq 7$, and no degree $\ell \geq 15$ reaches 0.5 in more than 10% of seeds. Clouds thus leave the hemispheric dichotomy and the largest continents recoverable, and nothing at scales below ~ 2000 km.

Land-ocean detection. We compute $d' = (\mu_{\text{land}} - \mu_{\text{ocean}})/\sigma_{\text{pooled}}$ within $|\vartheta| \leq 60^\circ$, using the true land labels. The truth field itself gives $d' = 3.97$. Cloud-free TDI gives $d' = 3.58 \pm 0.03$, the white GLS fit 3.61 ± 0.02 , and B2 3.04 ± 0.005 ; at 55.8% realized cover TDI gives $d' = 0.732 \pm 0.081$ and B2 0.674 ± 0.131 . Significance is bounded by the physical cloud-only null — 40 realizations in which the surface is replaced by its mean and only weather is retained — which has $\bar{d}' = 0.001$, $\text{sd} = 0.090$ and 95th percentile 0.114; the cloudy TDI value exceeds every null member, but a 40-member ensemble cannot support a p value below $1/(N_{\text{null}} + 1) = 0.024$. We therefore report $p \leq 0.024$ (rank-based, one-sided) for the cloudy detection rather than the v2 “ $p < 0.02$ against 100 surrogate nulls,” whose surrogate was a phase-scrambled equirectangular array — a geometry transform, not a photometric null. A label-permutation null on the same reconstructions is reported with the robustness suite in Section 5.7.

Quality is also a function of latitude, and the two effects do not point the same way. Scoring the archived fiducial reconstructions band by band over $|\vartheta|$ in 30° steps (archive `latband.npz`, 10 seeds) pairs each band with the illumination it actually receives — the campaign-mean Lambert factor over its disk pixels, and the campaign-mean column energy $\|F_{:,i}\|^2$, which is the information the operator delivers there.

Table 5 shows the ordering the referee expected and the one they did not. Expected: band quality is monotone in the information the operator actually delivers there. Across the well-sampled globe the two bands carrying 21.4 of column energy and 982 disk pixels each score 0.353 and 0.390 (cloudy) and 0.9814 and 0.9863 (cloud-free), while the two bands carrying 6.6 score 0.300 and 0.335, and 0.9655 and 0.9742 — the same ordering in both campaigns, with no band out of place. Unexpected: the two polar bands score highest in the cloudy case (0.451 and 0.547). That is an artefact of the metric, not a recovery claim, and the table gives the two diagnostics that expose it. The polar bands contain a half-illuminated ice cap against a dark ocean, so




Figure 4: Reconstruction fidelity vs. realized cloud cover across 10 paired seeds; bars are ± 1 standard error. The 16-bin phase-binned coadd is competitive below $\sim 40\%$ cover and falls behind as weather dominates (Table 4). Clouds, not rotation, control image quality at fixed observing resources.

Table 5: Latitude- and illumination-resolved quality for the fiducial campaigns (10 seeds, ± 1 SE). “Dayside” is the campaign-mean Lambert factor over the band’s disk pixels; “col. energy” is the campaign-mean column energy $\|F_{\cdot,i}\|^2$ of the operator; “dead” is the fraction of the band’s 432 grid cells whose column energy is below 10^{-3} of the maximum; “gain” is $\text{sd}(\text{recon})/\text{sd}(\text{truth})$ in the band, which Pearson r is blind to.

Latitude	disk pix	dayside	col. energy	dead	r (cloudy)	r (cloud-free)	gain (cloudy)
–90––60°	86	0.258	0.33	0.500	0.451 ± 0.038	0.9335 ± 0.0027	1.95
–60––30°	546	0.524	6.61	0	0.300 ± 0.020	0.9655 ± 0.0008	3.89
–30–0°	982	0.691	21.43	0	0.353 ± 0.015	0.9814 ± 0.0002	3.49
0–30°	982	0.691	21.43	0	0.390 ± 0.025	0.9863 ± 0.0001	2.77
30–60°	546	0.524	6.60	0	0.335 ± 0.022	0.9742 ± 0.0006	2.89
60–90°	86	0.258	0.33	0.486	0.547 ± 0.038	0.9358 ± 0.0025	1.68

their truth standard deviation is the largest of any band (0.148–0.160 against 0.097–0.137 mid-latitude) while their structure is almost entirely the low-degree cap edge: 65% and 71% of the two polar bands’ variance survives 1.5-pixel smoothing, against 71–82% for the four well-sampled bands. A reconstruction that gets nothing but the smooth cap still correlates well with truth. The amplitude column confirms it — the polar gain is 1.68–1.95 whereas the mid-latitude bands are inflated 2.8–3.9 $\times$ — and restricting the polar correlation to the cells the operator actually samples raises it further, to 0.518 and 0.652. Half of each polar band is dead: the circular $n = 64$ raster never samples $|\vartheta| \gtrsim 85^\circ$ (the 144 exactly-zero basis columns of Section 4.4), so those cells are set by the regularizer alone and contribute truth-like structure to nothing but the denominator of r . Cloud-free the same geometry gives the clean illumination monotonic trend, 0.9863 at the equator to 0.934 at the poles, with the polar deficit already concentrated in the dead cells (0.978–0.974 restricted to sampled cells). The band offsets

Table 6: Estimator-component ablation at $f_c = 0.55$ (10 paired seeds; means $\pm$ standard error). “Profiling” is the exact anchored nuisance elimination of Section 4.2; “deflation” is the v2 approximation to it; “whitening” is the OU temporal covariance.

Configuration	Pearson r	SSIM	NRMSE
A: white, no slot handling	0.2886 ± 0.0097	0.0853 ± 0.0031	0.966
B: OU whitened, no slot handling	0.2863 ± 0.0090	0.0825 ± 0.0031	0.974
C: v2 approximate deflation, OU	0.3422 ± 0.0082	0.1058 ± 0.0055	0.814
D: exact profiling, white, anchored	0.3454 ± 0.0082	0.1095 ± 0.0051	0.809
E: exact profiling, white, free	0.3453 ± 0.0082	0.1093 ± 0.0054	0.809
F: full TDI (exact profiling + OU)	0.3422 ± 0.0082	0.1060 ± 0.0053	0.814
F–A (Δr)	$+0.0536 \pm 0.0061$ ($t = 8.8$, 10/10)		$p = 1.1 \times 10^{-5}$
D–A (Δr)	$+0.0568 \pm 0.0046$ ($t = 12.4$, 10/10)		$p = 5.9 \times 10^{-7}$
D–C (Δr)	$+0.0032 \pm 0.0019$ ($t = 1.7$, 6/10)		$p = 0.13$
F–D (Δr)	-0.0032 ± 0.0019 ($t = -1.7$, 4/10)		$p = 0.13$

are symmetric-to-within-noise — $|\Delta a| \leq 0.019$ cloudy and ≤ 0.005 cloud-free — so the latitude dependence is a dependence of *correlation*, not a latitudinal albedo bias. We quote the band-resolved r as a coverage diagnostic and explicitly not as a claim of polar imaging.

5.4 Component-wise ablation of the inversion pipeline

Table 6 ablates the estimator itself across 10 paired seeds at nominal $f_c = 0.55$ (0.558 realized cover), holding the data, the operator, the grid and the regularizer fixed and changing only how the dwell-slot nuisance and the temporal covariance are handled. The v2 “ablation” varied forward-model choices and deflation heuristics together, so it could not assign credit; the four v2 rows also do not correspond to any subset of these six configurations.

Three conclusions follow, and they differ from v2.

(i) The gain is nuisance handling, and it is real. Exact profiling lifts r by $+0.054$ to $+0.059$ over the no-slot-handling fits (10/10 seeds), with $\Delta\text{SSIM} = +0.021 \pm 0.004$ ($t = 5.6$). The v2 attribution of the same effect to “slot deflation” with $\Delta r = +0.041$ and $p < 10^{-3}$ mixed two things: the correct idea (eliminate the per-slot mean) with an approximation that also rescales the marginal noise by $\sqrt{1 - 1/N_{\text{sc}}}$.

(ii) Exact profiling is not the same map as the approximation, but the fidelity difference is not resolved at 10 seeds. The relative L_2 difference between C and D is $9.3\% \pm 0.1\%$ of the map norm, i.e. the approximation changes the reconstruction materially, yet its effect on r ($+0.0032 \pm 0.0019$, $p = 0.13$) is within our power. With 10 paired seeds the distribution-free sign-test floor is $p = 2/2^{10} = 2 \times 10^{-3}$, so “not resolved” here means $|\Delta r| \lesssim 0.007$ at 95% confidence, not “equivalent”. We report the exact treatment because it is exact, not because it demonstrably images better.

(iii) Temporal whitening contributes essentially nothing at the fiducial cadence. $F-D = -0.0032 \pm 0.0019$ and $B-A = -0.0023$: with 16 revisits spaced 5.47 d and $\tau_{\text{cloud}} = 4$ d, consecutive passes are already nearly independent ($\rho_1 = 0.255$; Table 8), so the whitening gains little correlation and pays a small variance penalty. The term “time-domain inversion” therefore credits time registration and nuisance elimination; the OU covariance is defensible as a principled weighting, not as the source of the gain.

Finally, the anchored (D) and free (E) profiles give maps differing by $0.98\% \pm 0.27\%$ in norm and identical r to four decimals, while their zero-mode conditioning differs by a factor of 41 (Section 4.4). The anchor is what makes the *absolute* albedo scale identifiable; it does not improve correlation. Any claim about retrieved albedo without the anchor is unsupported.

5.5 Isolating the causal mechanism of the cadence law

Figure 5 and Table 7 report the cadence law under the two resource ledgers of Section 3.4, at nominal $f_c = 0.55$. Each cell is a mean $\pm$ standard error over the six paired seeds of the campaign study; the four estimator columns are the profiled TDI, the white-noise GLS fit, the 16-bin phase-binned coadd, and the phase-blind coadd. The v2 table quoted r_{inst} , r_{smear} , r_{cf} and r_{iid} columns that the shipped pipeline never produced; they have been replaced by controls that are in the archive (Section 5.6).

5.6 What the cadence law does and does not say

The trend is monotone, and it is dominated by weather sampling rather than by conditioning. In arm A, r rises $0.049 \rightarrow 0.571$ as M_p goes $4 \rightarrow 64$, and every step is unanimous across the six paired seeds ($t = 8.7$ to 21.9 ; marginal gaps $+0.124, +0.173, +0.151, +0.075$). The gain saturates: the last doubling buys less than half of the previous one. The cloud-free control of Section 5.1 shows that the collapse

Table 7: Cadence law per resource arm (nominal $f_c = 0.55$, 6 paired seeds). Arm A holds $M_{p,t_s} = 8$ h and lets the wall clock exceed 90 d; arm B holds 89.87 d and lets t_s fall. η is the duty cycle under 45 s/dwell overhead.

Arm	M_p	t_s	η	wall	r_{TDI}	r_{white}	r_{B2}	r_{B1}
A	4	7200	0.994	85.73	0.049 ± 0.005	0.024 ± 0.003	0.063 ± 0.011	0.003 ± 0.020
A	8	3600	0.988	86.27	0.173 ± 0.013	0.144 ± 0.011	0.068 ± 0.011	0.063 ± 0.019
A	16	1800	0.976	87.33	0.346 ± 0.010	0.291 ± 0.009	0.321 ± 0.019	0.031 ± 0.010
A	32	900	0.952	89.47	0.497 ± 0.014	0.475 ± 0.012	0.074 ± 0.004	-0.019 ± 0.003
A	64	450	0.909	93.73	0.571 ± 0.010	0.573 ± 0.013	0.503 ± 0.013	0.179 ± 0.002
B	4	7548.8	0.994	89.87	0.049 ± 0.006	0.031 ± 0.005	-0.130 ± 0.005	-0.064 ± 0.017
B	8	3751.9	0.988	89.87	0.164 ± 0.014	0.140 ± 0.013	-0.134 ± 0.021	-0.024 ± 0.019
B	16	1853.4	0.976	89.87	0.325 ± 0.017	0.267 ± 0.015	-0.231 ± 0.022	0.015 ± 0.014
B	32	904.2	0.953	89.87	0.498 ± 0.008	0.489 ± 0.010	-0.140 ± 0.058	-0.025 ± 0.006
B	64	429.6	0.905	89.87	0.576 ± 0.012	0.624 ± 0.008	0.454 ± 0.018	0.076 ± 0.002

from 0.979 to 0.346 at the fiducial cadence is entirely the weather term, and more passes help because each pass is a fresh draw of it.

The two controls, now actually measured. The cloud-free and per-pass-independent controls are computed on the same campaign objects, seeds and estimator as the main sweep (six paired seeds, both arms). Cloud-free r : 0.773 ± 0.008 , 0.934 ± 0.001 , and then 0.979, 0.984, 0.987 to within ± 0.0002 for $M_p = 16, 32, 64$ in arm A, with arm B within 0.04 of each. Per-pass independent weather: 0.051, 0.165, 0.342, 0.494, 0.583, against the correlated-OU values 0.049, 0.173, 0.346, 0.497, 0.571 in the same order. Two conclusions follow, and they point in opposite directions. (i) *Conditioning contributes, but it saturates:* the cloud-free curve still rises with M_p , by 0.214 over $4 \rightarrow 64$, so part of the v2-era trend is coverage and not weather averaging. It is, however, a minority of the effect and it is front-loaded: it gains 0.16 in the first doubling ($M_p = 4 \rightarrow 8$, where a raster position receives only four dwells) and only 0.008 from $M_p = 16 \rightarrow 64$, while the cloudy curve is still climbing steeply there ($0.346 \rightarrow 0.571$, $+0.225$). (ii) *The law does not need OU structure:* the iid and OU reconstructions agree at every cadence — paired differences $+0.002$, -0.008 , -0.004 , -0.003 , $+0.012$ in arm A and $+0.001$, -0.010 , $+0.014$, -0.026 , $+0.002$ in arm B, every $p_i \geq 0.14$, none approaching the exact sign-flip floor of 0.031 attainable with six pairs. What revisits buy is independent samples of the atmosphere, whatever its temporal correlation function turns out to be; the OU model determines *how many* of the M_p looks are effectively independent (Table 8), it is not the reason that more looks help. The absolute cost of weather at fixed cadence, $\Delta r_{\text{cf-OU}}$, falls monotonically from $+0.724 \pm 0.010$ at $M_p = 4$ to $+0.416 \pm 0.011$ at $M_p = 64$ (arm A; both arms significant at $p_i < 4 \times 10^{-7}$), which is the same statement in the other direction: the cadence law is about averaging the weather down, not about the weather getting worse.

The two arms are statistically indistinguishable. Pairing arm A against arm B at matched M_p gives Δr of $-0.001, +0.008, +0.021, -0.002, -0.005$ with paired standard errors $0.005-0.018$ and $p \geq 0.058$ in every case. The

wall-clock cost of doubling the revisit count is therefore paid out of the photon budget with no measurable loss over the range studied: the law is about *independent looks*, not about integration time per look, so long as the dwell stays long enough to beat the overhead (duty cycle $0.994 \rightarrow 0.905$ across the arms).

Whitening loses its edge at the fastest cadence. The profiled estimator beats the white fit at every $M_p \leq 32$ ($+0.025$ to $+0.058$) but at $M_p = 64$ the ordering reverses in both arms: -0.002 ± 0.009 in arm A and -0.047 ± 0.009 (0/6 seeds) in arm B. At 1.47 d spacing the OU correlation $\rho_1 = 0.69$ (Table 8) is real, but the estimator’s assumed σ_{cl} is calibrated at $t_s = 1800$ s, and the whitening then overweights the residual structure. The safe reading is that temporal whitening pays between roughly 8 and 32 revisits, and the fiducial 16 is not the optimum of the whitened fit.

The phase-binned baseline is partly a coverage artifact, and we can show it. B2 does not merely degrade in arm B, it goes strongly *negative* (-0.23 at $M_p = 16$), which is not a statement about TDI. The cause is geometric. Registering samples into 16 rotational phase bins does not leave each bin with a full raster: counting, per bin, how many of the 4096 image-plane pixels receive no sample at all gives an average of 256 empty pixels per bin at the fiducial configuration (arm A, $M_p = 16$), but 2048 at $M_p = 16$ in arm B and 3584 at $M_p = 4$ — the phase pattern of a boustrophedon raster does not tile the bins uniformly, and the two arms place different dwells at different phases. `reconstruct_phasebin` fills an empty pixel with the bin’s raster mean and then back-projects, so a coverage-starved bin contributes a smoothed, self-referential map, exactly the artifact flagged for the 8-bin variant in Section 4.6 but an order of magnitude worse. The B2 column is therefore reported for completeness and is *not* used to argue the cadence law; the comparison we stand behind is TDI against the white-noise GLS fit, which shares the operator and has no rasterization step.

Overheads bound the law from below, not above. With $T_{\text{oh}} = 45$ s fixed, η falls from 0.994 to 0.905; extrapolating the same ledger to $t_s = 100$ s costs 31% of the photons and to $t_s = 45$ s costs half of them, so the “ultrafast cadence” limit is where overheads bite. The v2 sentence claiming overhead losses “offset statistical gains” at $t_s \lesssim 100$ s was not measured — no configuration below 429.6 s is in the archive — and is deleted.

Table 8 makes the cadence law quantitative in the only unit that matters, the number of *statistically independent* weather samples. At the fiducial $\tau_{\text{cloud}} = 4$ d, $M_p = 64$ buys 12.1 effective looks, not 64 — a 5.3-fold saturation that is invisible in M_p alone. If the extrapolated weather decorrelation time is shorter ($\tau = 1$ d), the same hardware buys 40.3 looks and the cadence law steepens; if it is longer, the law is nearly flat beyond $M_p = 16$. The cadence requirement of Section 6 should therefore be quoted jointly with the assumed τ_{cloud} : as a number of revisits alone it is not a mission specification.

Table 8: Effective number of independent looks, $N_{\text{eff}} = M_p^2 / (M_p + 2 \sum_{k \geq 1} (M_p - k) \rho_k)$ with $\rho_k = e^{-k \Delta t / \tau_{\text{cloud}}}$ and inter-pass spacing $\Delta t = 86,400 / (M_p / 90)$ s. Values are identical for both arms to the third digit because spacing is set by the wall clock.

τ_{cloud}	$M_p = 4$	8	16	32	64
spacing (d)	21.47	10.80	5.47	2.80	1.47
ρ_1 (lag-1)	0.005	0.067	0.255	0.497	0.693
$N_{\text{eff}}, \tau = 1$ d	4.00	8.00	15.87	28.44	40.31
$N_{\text{eff}}, \tau = 2$ d	4.00	7.94	14.16	19.66	22.92
$N_{\text{eff}}, \tau = 4$ d	3.97	7.11	9.83	11.23	12.11

Table 9: Spin-geometry mis-specification, nominal $f_c = 0.55$, 10 paired seeds $\pm$ standard error.

Mis-specification	Pearson r	SSIM	contrast
none (fiducial)	0.3422 ± 0.0082	0.1060	1.071
pole tilt 5°	0.2856 ± 0.0138	0.0828	0.895
pole tilt 20°	0.1699 ± 0.0177	0.0350	0.569
phase error 5°	0.2949 ± 0.0083	0.0811	0.929
phase error 15°	0.1874 ± 0.0085	0.0324	0.596
phase error 30°	0.0719 ± 0.0102	-0.0027	0.250
phase error 60°	-0.0269 ± 0.0138	-0.0064	-0.066

5.7 Broader robustness: spin geometry, clouds, and systematics

5.8 Spin-geometry ephemeris requirements

Table 9 replaces the v2 numbers with 10 paired seeds and the corrected operator. The tolerance is asymmetric and steeper than v2 claimed: a 5° pole tilt costs 0.057 in r ($0.3422 \rightarrow 0.2856$, i.e. a 17% relative loss, not “virtually intact”), and 20° costs 0.172 (50%). Phase error is worse still per degree: 5° costs 0.047, 15° costs 0.155, 30° leaves $r = 0.072$, and at 60° the reconstruction is uncorrelated with the truth ($r = -0.027 \pm 0.014$), i.e. the map is not degraded, it is about the wrong hemisphere. The land–ocean contrast ratio tracks this collapse monotonically ($1.07 \rightarrow 0.90 \rightarrow 0.57 \rightarrow 0.25 \rightarrow -0.07$).

The precursor requirement that follows is stated as a budget, not as a threshold: at 5° in either angle the campaign loses about one sixth of its correlation, so *both* must be held well inside 5° for the cloud-limited number of Section 5.2 to be the relevant one. v2’s “pole tilt $\lesssim 5^\circ$, initial phase $\lesssim 5^\circ$, period to $\sim 10^{-4}$ ” was numerically close but was never a measured tolerance; the measured one is in Table 9.

5.9 Is the spin state refinable in flight?

No, and v2’s claim that it was is withdrawn. Figure 6b is a 5×5 grid over assumed period error $\Delta P / P \in$




Figure 5: Cadence law under the two resource ledgers, all values read directly from the shipped archives. (a) Pearson r vs. registered revisits M_p for the profiled TDI, the white-noise GLS and the two coadd baselines; solid lines are arm A ($M_{p,t_s} = 8$ h, 6 paired seeds) and faded lines arm B (89.87 d wall clock). The trend is monotone and saturating, and the two arms agree within the paired standard error at every M_p . (b) The quantity that actually saturates: the effective independent look count N_{eff} at $\tau_{\text{cloud}} = 1, 2, 4$ d against the $N_{\text{eff}} = M_p$ line of perfectly decorrelated weather (Table 8), with the arm-A duty cycle under the 45 s/dwell overhead ledger on the right axis.

$\{-2, -1, 0, +1, +2\} \times 10^{-4}$ and assumed initial phase error $\Delta\phi_0 \in \{-10, -5, 0, +5, +10\}^\circ$, with the whitened χ^2 of the anchored and free profiles and the resulting r , on four seeds.

The objective contains no information about the phase. At fixed period, the whitened χ^2 changes by less than 10^{-11} in relative value across the whole $\pm 10^\circ$ phase range — not small, but numerically identical. The reason is structural: the surface basis is complete in longitude, so shifting the assumed initial phase is a relabeling of ϕ , and the minimizer follows it. A χ^2 landscape cannot refine a quantity that leaves χ^2 invariant, so the v2 sentence “residual misalignments are refinable in-flight via profile grid search” is false as written and has been deleted.

The period direction is monotone across the grid, not convex with a minimum at truth: $\bar{\chi}^2$ decreases from 42,747 at $\Delta P/P = -2 \times 10^{-4}$ to 42,493 at $+2 \times 10^{-4}$ for every seed, so the grid minimum sits on the boundary and no curvature estimate exists; a parabolic fit per seed returns vertices scattered over $(-3 \times 10^{-3}, +5 \times 10^{-4})$. Over this grid the anchored and free profiles differ by a constant 19.5 in χ^2 and agree on every argmin, so the flatness is not an artifact of the anchor.

What is informative is reconstruction quality: r peaks at $\Delta\phi_0 = 0$ and $\Delta P/P = 0$ for 4/4 seeds, and a 10° phase error costs 0.12 in r while a 2×10^{-4} period error costs 0.03. That is the right statement for a mission design: the spin state is *observable through image quality*, so it can be validated by checking that a map looks like a planet, but it is not recoverable by minimizing the photometric residuals, and the range over which period error is benign ($\lesssim 10^{-4}$, ~ 1 s over 86,400 s) is far narrower than the phase tolerance in rela-

tive terms. A full two-dimensional spin-state solution from these data alone is not supported by this experiment and is left open.

5.10 Cloud-climatology mis-specification

Correlation is affine-blind, albedo is not. Sweeping the assumed mean cover by $\Delta f_c \in [-0.15, +0.15]$ and the assumed cloud albedo by $\Delta A_{\text{cl}} \in [-0.15, +0.15]$ leaves Pearson r *exactly* constant at 0.3422 over the entire 5×5 grid, and SSIM, NRMSE and contrast are likewise functions of Δf_c alone. This is not a numerical accident to be quoted as robustness — it is the definition of r being invariant to affine transforms of its argument — and v2’s phrasing invited exactly that reading. The quantities that do move are the absolute ones: mean albedo bias runs from +0.206 at $(\Delta f_c, \Delta A_{\text{cl}}) = (-0.15, -0.15)$ to -0.551 at $(+0.15, +0.15)$, RMSE is minimized on the $\Delta A_{\text{cl}} > 0$ side of the grid (best value 0.616, worst 1.813), and the land–ocean contrast ratio is set by Δf_c alone, from 0.80 to 1.61. The honest requirement is therefore: cloud climatology to $\sim 10\%$ in cover for absolute albedo at the 0.1 level, and any *feature-detection* claim must not be quoted with an albedo scale it cannot support.

Structured weather. Three climatologies beyond the homogeneous OU field, each scored with the σ_{cl} calibrated for the homogeneous case, at matched realized cover (0.549–0.558): latitudinal ITCZ banding gives $r = 0.3240 \pm 0.0126$, a three-timescale field ($\tau = 1, 10$ d, weight 0.7) 0.3174 ± 0.0088 , and per-pass independent clouds 0.3151 ± 0.0093 . Against the homogeneous fiducial 0.3422 ± 0.0082 on the *same* seeds the paired degradations are -0.018 ± 0.019




Figure 6: **Robustness suite.** (a) Correlation vs. spin-pole tilt and assumed initial-phase error (10 paired seeds; the zero-error point is the fiducial fit). (b) Whitenened χ^2 , re-referenced to its own value at zero phase, as a function of assumed phase offset for every trial period on the grid: the relative variation is $\leq 10^{-11}$, i.e. the objective is exactly flat in phase, while the reconstruction quality r (right axis, green) peaks at the truth — the phase error is observable through image quality, not through the likelihood (Section 5.9). (c) Scale-dependent correlation $r(\ell)$ for the cloud-free and cloudy campaigns (10 seeds, ± 1 SE), with the $\ell \leq 20$ grid ceiling marked.

(ITCZ, $t = -1.0$), -0.025 ± 0.009 (multi- τ , $t = -2.8$) and -0.027 ± 0.016 (iid, $t = -1.7$): at this seed count none of the three is more than a modest, marginally resolved penalty, so the fiducial number is not an artifact of an over-smooth weather model. Note in passing that per-pass independent weather scores slightly *lower*, not higher, than correlated weather under the OU-modelled estimator — the cadence gain of Section 5.5 is about the number of looks, not about whether the field is formally decorrelated. Surface-coupled clouds ($\beta_{\text{surf}} = 0.25$: cloudiness enhanced over the dark ocean) are the exception, with $r = 0.1444 \pm 0.0124$, a paired loss of -0.198 ± 0.017 ($t = -11.3$, 0/10 seeds), and a contrast ratio of only 0.379. The failure mode is physical and important: a cloud pattern locked to the surface does not average out with revisits, because it is phase-correlated with the very signal it is meant to contaminate. This is the one climatological structure the cadence law cannot buy off, and it is the strongest argument in the paper for treating surface-cloud coupling as a systematic rather than as noise.

5.11 Surface morphology: fragmentation is free, concentration is not

The referee asked whether the result depends on “substantially different surface-map morphologies.” v2 asserted robustness to morphology; that scan was produced with the superseded quadrature and the approximate deflation estimator, so it is re-measured here rather than inherited. Three truth fields are built from the same seed and the same quantile land fraction — 30% of the fine grid above the albedo threshold, giving mean albedo 0.235–0.237 and standard deviation within 2% across the three — differing only in how the land is *arranged*: the fiducial *earthlike* field (fragmented continents), a *supercontinent* field whose

Table 10: Surface-morphology sensitivity, 10 seeds each, paired against the earthlike reconstruction on the same seeds. The cloud-free columns separate a resolution effect from a weather effect: all three morphologies are recovered to $r \geq 0.974$ without clouds, so the cloudy spread is a weather-surface coupling, not an information loss.

Truth field	σ_{cl}	$r, f_c = 0$	$r, f_c = 0.55$	paired Δr vs. earthlike	d' (cloudy)
earthlike (fiducial)	0.4406	0.9786 ± 0.0002	0.3422 ± 0.0082	—	0.732 ± 0.081
supercontinent	0.4870	0.9844 ± 0.0002	0.3026 ± 0.0111	-0.0396 ± 0.0090	0.725 ± 0.036
archipelago	0.4426	0.9738 ± 0.0002	0.3393 ± 0.0078	-0.0030 ± 0.0056	0.728 ± 0.019

power is pushed to the largest scales (a Pangaea-like single mass, 48% of coarse pixels land-bearing against the fiducial 51%), and an *archipelago* field whose power is pushed to the smallest scales (50%). Each gets its own weather realization on the same cloud seeds and its own calibrated σ_{cl} , because the cloud-induced variance depends on the albedo pattern it modulates: $\sigma_{\text{cl}} = 0.4406$ (earthlike), 0.4870 (supercontinent), 0.4426 (archipelago).

Table 10 answers the question in two parts, and the honest answer is that the v2 claim was half right. Breaking the continents into island chains costs essentially nothing: 0.3393 ± 0.0078 against the fiducial 0.3422 ± 0.0082 , a paired difference of -0.0030 ± 0.0056 , smaller than one standard error and with the sign split 4/10 seeds. Fragmentation at $\ell \gtrsim 9$ is a scale the cloudy campaign has already lost (Section 5.3), so a finer truth map is not a harder inverse problem. Concentrating the same land into one large mass does cost something: -0.0396 ± 0.0090 ($p_t = 0.0017$, though the sign is negative in 8/10 seeds and the exact sign-flip p is 0.088). That is 12% of the fiducial value: one fifth of the surface-coupling penalty of Section 5.10, and half again the multi-timescale effect, so it belongs in the same tier of mis-

Table 11: Weather-parameter mis-specification. The true field is varied in correlation length, advection rate and decorrelation time at matched realized cover; $\sigma_{\text{cl}}^{\text{var}}$ is the cloud amplitude calibrated for that field (fiducial 0.4406). “Recalibrated – misspecified” is the paired difference over the 10 seeds with its standard error and exact sign-flip p .

True field	$\sigma_{\text{cl}}^{\text{var}}$	r , misspec.	r , recal.	paired Δr	p_{sign}
corr. 8°	0.4158	0.3163 ± 0.0101	0.3163 ± 0.0101	$+1.3 \times 10^{-6}$	0.088
corr. 20°	0.4897	0.3269 ± 0.0142	0.3269 ± 0.0142	-2.0×10^{-6}	0.088
adv. 3°/d	0.4475	0.3133 ± 0.0097	0.3133 ± 0.0097	-3.2×10^{-7}	0.088
adv. 12°/d	0.4389	0.3237 ± 0.0130	0.3237 ± 0.0130	$+8.1 \times 10^{-8}$	0.088
$\tau = 2$ d	0.4416	0.3160 ± 0.0074	0.3176 ± 0.0072	$+0.0017 \pm 0.0008$	0.088
$\tau = 8$ d	0.4387	0.3259 ± 0.0133	0.3188 ± 0.0136	-0.0072 ± 0.0016	0.020

specification costs rather than in the benign one. The mechanism is the one Section 5.10 identifies rather than a loss of information: the supercontinent’s σ_{cl} is 10.5% larger than the fiducial’s, because a single wide landmass presents long, coherent albedo edges for the weather to modulate, and its contrast collapses from 1.071 to 0.983 — the smoothed solution fills in the ocean around the continent instead of sharpening its coast. Cloud-free the same three fields score 0.9738–0.9844, all within 0.006 of each other and with the *supercontinent highest* (+0.0058, 10/10 seeds), which is what one expects when there is no weather and the large-scale structure is easy: the morphology penalty exists only in the presence of clouds.

We therefore replace the v2 blanket assertion with a bounded one: the pipeline is insensitive to surface *granularity* at fixed land fraction, and mildly sensitive to surface *coherence* at fixed land fraction, through the cloud amplitude the surface generates rather than through the inversion itself. A single-continent planet is a ~ 0.04 problem in r , not a qualitative failure, and the direction of the effect — concentrated land is worse, because it makes the weather more visible — is a prediction the surface–cloud coupling section makes and this table tests.

5.12 Weather parameters, and the amplitude the estimator never uses

The referee asked for robustness to “at least two additional spatial correlation lengths, advection rates and decorrelation times.” Table 11 gives six such fields — correlation length 8° and 20° (fiducial 12°), advection 3° day^{−1} and 12° day^{−1} (fiducial 6°), and decorrelation time 2 d and 8 d (fiducial 4 d) — each a different true cloud field at matched realized cover, reconstructed on the same 10 seeds and scored two ways: *misspecified* keeps the fiducial OU τ and the σ_{cl} calibrated for it, *recalibrated* substitutes the variant’s own τ and σ_{cl} .

Two things stand out, and the second was not expected. Across the four fields that change the spatial or temporal *structure* of the weather, r moves by 0.313–0.327 — within one another’s error bars, and within ~ 0.02 of the fidu-

cial 0.3422 ± 0.0082 — and changing the assumed climatology to the true one buys nothing at all: the paired differences are at the 10^{-6} level, far below the resolution of 10 seeds but also far below anything worth a mission requirement. The reason is the second finding, which we tested directly. Figure 7c scores the fiducial cloudy reconstruction as a function of the *assumed* σ_{cl} , from 0 (weather treated as uncorrelated) to $16\times$ the calibrated value: the mean paired change in r never exceeds 2.6×10^{-4} anywhere between $0.25\times$ and $16\times$ (no individual seed moves by more than 5.2×10^{-4}), and even with the cloud term removed entirely it is $3.2 \times 10^{-3} \pm 1.9 \times 10^{-3}$ (0.3454 versus 0.3422, only 6/10 seeds) — while the whitened objective over the same range spans a factor of 8×10^4 , from $3.2 \times 10^2 \times$ the fiducial value at $\sigma_{\text{cl}} = 0$ to $3.9 \times 10^{-3} \times$ at $16\times$. *The reconstruction is essentially insensitive to the cloud amplitude it is told to expect, and the χ^2 is essentially only a function of it.* The insensitivity is not an artefact of the anchoring: running the same sweep on the two estimators this paper compares TDI against — the free, un-anchored profiler and v2’s approximate deflation — gives per-seed spreads across $0.25\times$ – $16\times$ of $2.74 \times 10^{-4} \pm 0.63 \times 10^{-4}$ for the anchored profiler, $2.72 \times 10^{-4} \pm 0.63 \times 10^{-4}$ for the free profiler and $2.56 \times 10^{-4} \pm 0.59 \times 10^{-4}$ for v2 deflation (archives `sigma_sensitivity_free`, `sigma_sensitivity_v2`); all three are flat at the 10^{-4} level against a seed-to-seed scatter of 0.026. This is why “recalibrated” and “misspecified” are indistinguishable in Table 11, and it sharpens rather than softens the referee’s concern: σ_{cl} enters the pipeline through the whitening and the slot-covariance blocks, whose contribution to the solution is dominated by the operator’s own time-domain structure, so an atmospheric climatology that is wrong by a factor of four is not a fidelity hazard, whereas surface–cloud coupling (Section 5.10) — which changes the *structure*, not the amplitude — costs 0.198. The practical consequence for mission design is that the pre-flight requirement is on weather *correlation structure*, and σ_{cl} itself can be fit in flight from the residuals even crudely, since χ^2 does respond to it steeply even though r does not.

5.13 Regularization: the frozen weight is not optimal, and the objective cannot choose it

The v2 pipeline fixed $\lambda = 3 \times 10^{-3}$ on a held-out calibration seed and never swept it (minor comment 3). Fig. 7a,b is that sweep, over six decades (10^{-5} – 10) on 10 seeds, for the cloud-free and fiducial cloudy campaigns, scored by paired difference against the frozen value and by the whitened objective.

The answer to the referee’s question is that the reconstruction is *not* insensitive to λ , and the frozen value is not the optimum. Cloud-free, r rises monotonically from 0.9752 at $\lambda = 10^{-5}$ to a maximum of 0.9916 at $\lambda = 0.3$ (+0.0130, 10/10 seeds, $t = 66.9$), turns over, and falls to 0.9639 by $\lambda = 10$; SSIM traces the same curve ($0.839 \rightarrow 0.945 \rightarrow 0.891$).




Figure 7: **Regularization and assumed-cloud-amplitude sensitivity (10 seeds).** (a) Paired Δr against the frozen $\lambda = 3 \times 10^{-3}$: cloud-free quality peaks near $\lambda = 0.3$ and then declines, while cloudy quality keeps rising to the largest weight tested (+0.178 at $\lambda = 10$). Bars are ± 1 SE. (b) The whitened objective $\chi^2/\chi^2(3 \times 10^{-3})$ cannot make the same choice: it is flat to $\sim 0.1\%$ over five decades in the cloudy case and rises only past $\lambda \sim 0.1$ in the cloud-free case. (c) The *assumed* cloud amplitude: the mean paired change in r is $\leq 2.6 \times 10^{-4}$ from $0.25\times$ to $16\times$ the calibrated σ_{cl} (and within 0.003 even at zero), while χ^2 spans a factor 8×10^4 .

Cloudy, r rises *monotonically across the entire tested range*, $0.3378 \rightarrow 0.5200$, a paired gain of $+0.178 \pm 0.007$ at $\lambda = 10$ (10/10 seeds), which is more than three times the entire algorithmic gain of the pipeline (Section 5.4). Meanwhile the objective is useless as a selector: in the cloudy case χ^2 varies by less than 0.1% from $\lambda = 10^{-5}$ to 0.3 — within ± 54 absolute, against an expected sampling width of $\sqrt{2N} = 362$ for $N = 65,536$ samples — and only at $\lambda \geq 1$ does it begin to rise (+193, +648, +1394). There is no discrepancy-principle minimum in the cloudy objective at all.

Because a monotone $r(\lambda)$ under a smoothing prior invites exactly the wrong reading — that more regularization is “better imaging” — we scored what the large- λ gain is made of. At four weights ($\lambda = 3 \times 10^{-3}$, 0.1, 1, 10) we recomputed the scale-dependent correlation of Section 5.3, the land-ocean d' and the contrast (4 seeds, archive `lambda_resolution.npz`). **Cloudy:** the harmonic bands are *unchanged* within the error — $r(\ell)$ over $\ell = 1-4$ stays at 0.72–0.74 and over $\ell \geq 13$ at 0.26–0.29 as λ goes $3 \times 10^{-3} \rightarrow 10$ — while d' rises from 0.749 ± 0.114 to 1.219 ± 0.191 and the contrast falls from 1.071 ± 0.121 to 0.929 ± 0.116 . The $\ell = 9-12$ band in fact slips slightly ($0.446 \rightarrow 0.381$), which is what smoothing a signal-starved band looks like. So the cloudy gain is real large-scale signal recovery, not a blur that trades resolution for correlation: clouds have already removed the $\ell \gtrsim 9$ information (Section 5.3) and λ cannot put it back, but at $\ell \leq 4$ — the hemisphere-scale land-ocean dichotomy the paper actually claims — extra smoothing recovers correlation. Cloud-free the control behaves as expected for a resolution claim: the $\ell \geq 13$ band degrades from 0.992 to 0.850, the $\ell = 9-12$ band from 0.998 to 0.854, and r turns over ($0.9787 \rightarrow 0.9906 \rightarrow 0.9639$) exactly where that loss overtakes the noise suppression.

This is a result we report rather than adopt. The headline numbers of this paper keep $\lambda = 3 \times 10^{-3}$, for three reasons: our scoring metric is correlation with a *known* truth map, so $r(\lambda)$ is available here and not in flight; the cloudy gain is a gain in a metric evaluated against a map whose power at $\ell \gtrsim 9$ is already gone, so it cannot be read as improved surface imaging; and the objective gives no independent way to select the weight from data alone, which is precisely the situation in which a frozen, truth-blind constant is defensible. What we no longer claim is that the choice is unimportant. A regularisation criterion that operates without truth — discrepancy principle, GCV, or a validated prior over surface models — is the one methodological gap this sweep exposes, and it is stated as such in Section 7.

5.14 Nulls, and what “significance” means here

For each of the 10 paired seeds the cloudy reconstruction gives $d' = 0.732 \pm 0.026$ against the true land labels. The label-permutation null (2000 permutations of the labels on the same map) gives $p = 5.0 \times 10^{-4}$ for every seed, which is its own floor ($1/(2000 + 1)$); the physical cloud-only null of Section 5.3 gives $p = 0.024$ for every seed, which is also its floor ($1/(40 + 1)$). Both are saturation statements, not measurements: what they show is that the observed separation exceeds every null draw, and the ensemble sizes bound the number we can honestly print. With 10 paired seeds the smallest attainable two-sided sign-test p is $2/2^{10} = 1.95 \times 10^{-3}$; no comparison in this paper is quoted below that value, and the v2 “ $p < 10^{-3}$ ” has been replaced by exact per-comparison values throughout.

5.15 Coronal and instrumental systematics: what the noise model actually tolerates

The reconstruction pipeline above assumes the corona contributes only Gaussian shot noise with a known mean. The referee asked whether that assumption is load-bearing, and v2 answered it by assertion. Here it is measured: the two non-Gaussian terms the referee names are injected into the fiducial cloudy campaign ($f_c = 0.558$, profiled estimator, GL-16 covariance) and the map is re-solved on each of the 10 seeds, whose unperturbed score is $r = 0.3422 \pm 0.0082$.

The units convention carries the entire result, so it is stated once, explicitly. The data vector is the normalized planet photocentre signal, hence a fractional residual ε of the coronal background enters the light curve as $\Delta y = \varepsilon Q_{\text{cor}}/Q_{\text{exo}} = 7.74 \times 10^4 \varepsilon$ in planet-reference units. The v2 injection applied its $\sim 10^{-4}$ amplitudes directly to y — a fractional error of the planet — which is $\sim 7.7 \times 10^4$ times too benign to constitute a test of a coronal residual at all. Two templates are swept separately and jointly: a coherent multiplicative gain ramp across the campaign, $1 + \varepsilon(2u - 1 + \frac{1}{2}(2u - 1)^2)$ with $u = t/t_{\text{end}}$, and a low-spatial-frequency streamer field of ~ 8 -cell transverse scale whose modes random-walk with correlation times of 15–60 d.

Table 12 is the answer. The tolerance is set near $\varepsilon \sim 10^{-6}$ of the coronal background: at 10^{-6} the paired loss is -0.0023 ± 0.0006 (drift, $t = -3.6$, $p = 6 \times 10^{-3}$), -0.0015 ± 0.0002 (streamer, $t = -6.3$, $p = 1.4 \times 10^{-4}$) and -0.0037 ± 0.0006 (both, $t = -5.9$, $p = 2 \times 10^{-4}$) — 0.4–1.1% of r , resolved as a real effect and negligible for every science claim in the paper. One decade later the field has lost 21% (drift), 23% (streamer) and 32% (both) of its correlation, and by $\varepsilon = 10^{-4}$ it has collapsed: $r = 0.084$, 0.036 and 0.057, with SSIM down to 1.3×10^{-3} , 7×10^{-4} and 7×10^{-4} , and the gain ramp driving the mean-albedo bias to +1.29 where the streamer drives it to -0.21 . The two terms do not add. The joint loss at 10^{-5} is -0.110 against a sum-of-parts -0.151 (73%), and at 10^{-4} -0.285 against -0.565 (50%); a coherent ramp is not decorrelated by revisits while the streamer partly is, so the joint curve is sub-additive without being milder in absolute terms, and it is the joint curve that a budget should be taken from.

What makes the tolerance a requirement rather than a number is the comparison scale. The single-dwell photon fluctuation, expressed in the same coronal units, is 2.99×10^{-7} : the $\varepsilon = 3 \times 10^{-7}$ row of Table 12 is therefore exactly one photon- σ of one dwell, and the reconstruction loses $\leq 0.15\%$ of r at it. The knee of the curve sits at 3–4 σ of that same quantity, not at some unmeasurably tiny fraction of the corona. Conversely, an independent background reference observing for one dwell ($t_b = \tau_{\text{int}} = 7200$ s) adds $\sqrt{Q_{\text{cor}}/(Q_{\text{exo}}^2 t_b)} = 0.02317$ of noise — equal to the photon σ to within 7×10^{-6} relative, as Section 4.5 derives — so a dedicated differencing pass buys the coronal mean and buys it at full photon cost, while leaving the low-order spatio-temporal structure that Table 12 prices entirely un-

Table 12: Measured effect of injected systematics on the fiducial cloudy reconstruction ($f_c = 0.558$, profiled estimator), averaged over the 10 seeds; the unperturbed value is $r = 0.3422 \pm 0.0082$. Δr is the paired difference against the same seed’s unperturbed fit, with t from a two-sided paired t -test ($n = 10$, so the sign-test floor is 1.95×10^{-3}). “Coronal units”: the residual is a fraction ε of the coronal background, so $\Delta y = 7.74 \times 10^4 \varepsilon$ in planet units. The $\varepsilon = 3 \times 10^{-7}$ row is marked because it is exactly 1σ of single-dwell photon noise expressed in the same units.

ε	Gain drift			Streamer			Both	
	r	Δr	t	r	Δr	t	r	Δr
0	0.3422	—	—	0.3422	—	—	0.3422	—
10^{-8}	0.3422	−0.0000	−1.8	0.3422	−0.0000	−1.9	0.3422	−0.0000
10^{-7}	0.3421	−0.0001	−2.0	0.3422	−0.0001	−2.3	0.3421	−0.0002
3×10^{-7} (1σ)	0.3418	−0.0005	−2.3	0.3420	−0.0002	−3.3	0.3416	−0.0007
10^{-6}	0.3400	−0.0023	−3.6	0.3408	−0.0015	−6.3	0.3385	−0.0037
10^{-5}	0.2700	−0.0723	−16.5	0.2640	−0.0782	−26.7	0.2319	−0.1103
10^{-4}	0.0842	−0.2581	−32.9	0.0355	−0.3067	−41.3	0.0568	−0.2854

touched. The mission requirement this section leaves behind is accordingly on the corona, not the planet: calibrate or model the low-order spatio-temporal modes of the background to the sub- 10^{-6} -of-background level, and carry a coherent gain ramp as a systematic with its own nuisance block rather than as noise.

The robustness claims of v2 are resolved here in three ways: restated with numbers, measured with a qualification, or withdrawn. The claim that cloud-induced noise inflation is “ ~ 14 – $23\times$ ” becomes the measured 12.8 – $21.0\times$ in σ , 163 – $440\times$ in variance, with the per-cover σ_{cl} values quoted outright (Section 5.3). The v2 assertion that the reconstruction is robust to surface morphology is now measured rather than inherited, and the honest answer is qualified: granularity costs nothing, coherence costs 0.040 in r (Section 5.11). The v2 frozen regularizer is likewise no longer asserted to be unimportant — a six-decade sweep shows the fiducial λ is far from optimal and that the objective cannot select it (Section 5.13), and the assumed cloud amplitude turns out to be one the estimator barely uses at all (Section 5.12). What is withdrawn without replacement is the v2 latitude-averaged “resolution” figure and the v2 surrogate-null p value, both re-derived in Section 5.3. Sensitivities re-measured with the corrected pipeline are reported in Sections 5.8, 5.9, 5.10, 5.11, 5.12 and 5.13. Where a v2 claim has not been re-run, this paper no longer makes it.

6 From Algorithm to Mission

6.1 Target selection: a nominal-proxy table, not an uncertainty budget

Table 13 lists the five nearest confirmed temperate planets around M dwarfs plus the solar-type radial-velocity candidate τ Cet e as SGL imaging targets. It is generated by `src/targets.py` and is explicitly a *nominal-proxy* table. The v2 version was labelled “propagated physical uncer-

tainties” and quoted $\pm$ values and confidence-style brackets; that wording is withdrawn, for three reasons, and what the table actually contains is stated below.

Radii. None of these planets transits, so the observed quantity is $M \sin i$, not M and not R_p . We use the Chen & Kipping (2017) *Terran* branch

$$R_p = 1.01 \left(\frac{M \sin i}{M_\oplus} \right)^{0.279} R_\oplus, \quad (14)$$

whose intrinsic dispersion is 4.03%, not the 15% quoted in v2 — 14.6% is the dispersion of the *Neptunian* branch, and 15% was attached to the wrong population. The $\pm$ column now carries the correct 4.03%, while the brackets use a deliberately generous $\pm 15\%$ radius allowance that is labelled as an assumption. The branch *transition* is at $2.04^{+0.66}_{-0.59} M_\oplus$; two rows (GJ 273 b, τ Cet e) lie above it and are flagged, because applying the rocky branch there is an extrapolation of the published relation rather than its stated domain of validity.

Inclination. v2 folded a mean de-inclination correction into the quoted radius, which is not the correct operation: the radius is a nonlinear function of the true mass, so the mean of the ratio differs from the ratio of the means. With $\cos i$ uniform on $[0, 1]$ and eq. 14, the conditional factor $R_p/R_\oplus(M \sin i) = (\sin i)^{-0.279}$ has mean

$$\langle (\sin i)^{-\alpha} \rangle = \int_0^1 (1-u^2)^{-\alpha/2} du = \frac{\sqrt{\pi} \Gamma(1-\alpha/2)}{2 \Gamma(3/2-\alpha/2)}, \quad (15)$$

which for $\alpha = 0.279$ is 1.0978 (the referee’s 1.0982 corresponds to $\alpha = 0.28$), with median 1.0409, 95th percentile 1.3837 and 99th percentile 1.7271. The often-quoted $(\langle \sin i \rangle)^{-\alpha} = (\pi/4)^{-\alpha} = 1.0697$ is a different number (ratio of means vs. mean of ratio) and is not used. These are geometric ratios conditional on a fixed minimum mass and a single power-law branch — *not* a posterior: a radial-velocity mass prior, detection biases and branch transitions all change them. Accordingly the nominal R_p in Table 13 carries no inclination correction at all, and the factors above are quoted separately.

Signal strength. The scalar proxy recommended by the referee is:

$$\text{SNR}_C(1800\text{s}) = 43.16 \left(\frac{A}{0.30} \right) \left(\frac{S}{S_\oplus} \right) \left(\frac{R_p}{R_\oplus} \right) \left(\frac{30\text{pc}}{z_0} \right), \quad (16)$$

which we state explicitly rather than leave as an implied scaling. It is a *bolometric scalar rule of thumb*, not a wavelength-integrated photon forecast; the latter requires the spectral SGL treatment (Turyshv & Toth 2022c, 2023b). Eq. 16 reproduces the referee’s check values to the last quoted digit for every row, and is compared there against the full optical calculation (SNR_C from the actual $F^\top C^{-1} F$ diagonal): Proxima Cen b 665.8 vs. 665.7, τ Cet e 909.2 vs. 909.1. Because eq. 16 is linear in A and R_p , the bracket in the table is the bounded product of $A \in$

$[0.15, 0.45]$ with the $\pm 15\%$ radius allowance; it is an assumption range, not a confidence interval.

Tracking cost. The image cylinder must be steered to follow the planet’s projected transverse acceleration. For a Keplerian orbit that quantity is not the commonly used circular estimate but

$$\Delta v(T) = \oint |\mathbf{a}_\perp(t)| dt, \quad (17)$$

with the projected transverse acceleration of a Keplerian orbit, $\mathbf{a}_\perp = -(GM_\star/r^2)(\cos E - e, \sin E \sqrt{1-e^2} \cos i)$, integrated in *time* over the window (the eccentric-anomaly measure $dM = (1 - e \cos E) dE$ supplies the weighting; averaging it away biases low- e rows and high- e rows differently). v2 quoted a $\pm$ “eccentricity modulation” error bar on Δv_{90} that could not be reproduced; it is replaced by the explicit range over the starting mean anomaly within the 90-day window and by the corresponding edge-on value. The circular limit $\Delta v = 4\pi^2 a T / P^2$ evaluated at z/z_0 is recovered as the $e \rightarrow 0$ case of eq. 17 to the quoted precision, and the edge-on/face-on ratio converges to the analytic $2/\pi = 0.6366$ (referee eq. R29); τ Cet e, with $P = 162.9$ d, needs only 0.11 km s^{-1} over 90 d, while the short-period Proxima b needs 5.81 km s^{-1} — a factor 53 spread that the v2 table compressed into an apparent $\sim 10\%$ uncertainty.

Tidal synchronization caveats for M-dwarf planets: While close-in M-dwarf planets are commonly presumed tidally locked ($P_{\text{rot}} = P_{\text{orb}}$), a precise orbital period does not by itself guarantee synchronous rotation. Non-zero orbital eccentricity drives pseudosynchronization to higher spin rates ($P_{\text{rot}} < P_{\text{orb}}$), capture into spin-orbit resonances (e.g., 3:2), and longitudinal librations, while atmospheric thermal tides can maintain asynchronous spin. Characterizing the true rotation period to $\sim 10^{-4}$ via precursor photometric light curves remains essential.

6.2 Deep-space transportation and swarm sizing

Cruising to $z = 650$ AU within 20 yr requires $v_\infty \simeq 154 \text{ km s}^{-1}$ (32.5 AU yr^{-1}), far exceeding chemical propulsion ($3\text{--}4 \text{ AU yr}^{-1}$; Turyshv 2026d). Deep-perihelion solar sailcraft deploying near $r_p = 0.05 \text{ AU}$ ($10.7 R_\odot$) achieve $v_\infty \simeq 105\text{--}155 \text{ km s}^{-1}$ at system areal densities $\sigma_{\text{tot}} \simeq 2.3\text{--}4.9 \text{ g m}^{-2}$. A 100 kg craft requires a $140 \times 140 \text{ m}$ sail operating near $\sim 740 \text{ K}$ (Helvajian et al. 2023; Turyshv 2026d).

Beyond 500 AU, solar flux is $> 4 \times 10^5$ times weaker than at 1 AU; all operational power must be nuclear. The Atomic Planar Power for Lightweight Exploration (APPLE) architecture—modular $10 \times 10 \times 1.7 \text{ cm}$ tiles integrating ^{238}Pu with thermoelectrics and solid-state batteries delivering $\geq 1.2 \text{ W}_e$ and 5 Wh storage per tile (NASA/Aerospace Corp. 2022)—scales to $\sim 10^2 \text{ W}_e$ at masses compatible with sailcraft.

Photometric science data volume is modest, and v2 overstated it by three orders of magnitude: the fiducial campaign is $65,536 \text{ samples} \times 32 \text{ bytes} = 2.10 \text{ MB}$ (the 32 bytes

Table 13: Nearest confirmed temperate planets (and one solar-type RV candidate) as SGL imaging targets – explicitly a *nominal-proxy* table, not a propagated uncertainty budget. Radii use the Chen & Kipping (2017) Terran branch $R = 1.01 (M \sin i / M_{\oplus})^{0.279} R_{\oplus}$ evaluated at $M \sin i$ (no de-inclination correction is folded into the nominal value; the conditional isotropic factors are given in the text); $\pm$ values are the 4.03% intrinsic dispersion of that branch. SNR_C follows the scalar rule of eq. 16 and the bracket is the bounded product of $A \in [0.15, 0.45]$ with a $\pm 15\%$ radius allowance – an assumption, not a confidence interval. Δv_{90} is the time-integral of the projected transverse acceleration over 90 d for a face-on orbit, with the range over starting mean anomaly and the corresponding edge-on value in parentheses. Observing geometry $z = 650 \text{ AU}$, $d = 1 \text{ m}$, $\lambda = 1 \mu\text{m}$, 1800 s .

Planet	Host (SpT)	z_0 (pc)	$M \sin i$ ($M_{\oplus}$)	$R_p \pm \text{disp.}^d$ ($R_{\oplus}$)	S ($S_{\oplus}$)	P (d)	D_{img} (km)	SNR_C [range] ^b (1800 s)	Focal direction ^c (J2000)	β_{ecI} (deg)	Δv_{90} (face) ^b (km s^{-1} [range], edge)
Proxima Cen b	Proxima Centauri (M5.5V)	1.30	1.07	1.03 ± 0.04	0.65	11.19	31.8	666 [283–1148]	$02^{\text{h}}30^{\text{m}}, +62^{\circ}41'$	+44.8	5.81 (3.70)
Ross 128 b	Ross 128 (M4V)	3.38	1.35	1.10 ± 0.04	1.38	9.87	13.1	581 [247–1003]	$23^{\text{h}}48^{\text{m}}, -00^{\circ}48'$	+0.5	2.94 (1.87)
GJ 1061 d	GJ 1061 (M5.5V)	3.67	1.64	1.16 ± 0.05	0.69	13.03	12.7	282 [120–487]	$15^{\text{h}}36^{\text{m}}, +44^{\circ}31'$	+60.8	1.68 (1.07)
Teegarden c	Teegarden’s Star (M7V)	3.83	1.11	1.04 ± 0.04	0.37	11.41	10.9	130 [55–224]	$14^{\text{h}}53^{\text{m}}, -16^{\circ}53'$	-0.3	1.72 (1.10)
GJ 273 b ^a	GJ 273 (M3.5V)	3.80	2.89	1.36 ± 0.06	1.06	18.65	14.4	490 [208–846]	$19^{\text{h}}27^{\text{m}}, -05^{\circ}14'$	+16.5	1.34 (0.85)
τ Cet e ^a	τ Ceti (G8.5V)	3.60	3.93	1.48 ± 0.06	1.71	162.90	16.5	909 [386–1568]	$13^{\text{h}}44^{\text{m}}, +15^{\circ}56'$	+24.8	0.11 (0.07)

^a $M \sin i$ lies above the Chen & Kipping (2017) Terran/Neptunian transition ($2.04^{+0.66}_{-0.59} M_{\oplus}$), so the rocky-branch radius used here is an extrapolation, not a measurement. ^b The SNR_C bracket is a *bounded range* over the stated albedo ($A \in [0.15, 0.45]$) and radius allowance ($\pm 15\%$), not a confidence level and not propagated from catalogue uncertainties. The Δv_{90} entries give the face-on value with the corresponding edge-on value in parentheses; the spread over the starting mean anomaly inside a 90-day window is below 0.02 km s^{-1} for every row except τ Cet e (0.09 – 0.13 km s^{-1}). ^c Focal direction is the antipodal celestial coordinate of the host; β_{ecI} is its ecliptic latitude, which sets the departure plane-change cost. ^d Radii, stellar parameters and orbital elements are from the references in Section 6.1; none of these planets transits, so $M \sin i$ (not M) is the observed quantity.

being time tag, value, uncertainty and bookkeeping), and the reconstructed surface at the paper’s own 36×72 resolution is 10.4 kB in float32. Even a 1024×1024 albedo map with per-pixel uncertainty is 8.4 MB, not “ ~ 2 GB for megapixel maps” — the v2 figure is off by $\sim 250\times$ and appears to have assumed an uncompressed float64 stack of every harmonic band. Total mission telemetry (navigation, star trackers, health, inter-craft ranging) will dominate the downlink by orders of magnitude, but laser relays hopping along the string-of-pearls swarm easily support tens to hundreds of kb s^{-1} (Helvajian et al. 2023), so science return is not bandwidth-limited.

Re-evaluating swarm extrapolations: The cadence law supports revisits improving fidelity, but the saturation is now measured rather than assumed: in arm A the four doublings from $M_p = 4$ to 64 buy $+0.124, +0.173, +0.151, +0.075$ in r , and at the assumed $\tau_{\text{cloud}} = 4 \text{ d}$ the effective-look count is 12.1 at $M_p = 64$ (Table 8), so the curve is already bending before the measured range ends. Extrapolating through several hundred visits to claim $r \geq 0.8$ is therefore an *asymptotic architectural scenario*, not a requirement, and it is now quantitatively bounded by N_{eff} : if the real weather decorrelation time is longer than 4 d, $M_p > 32$ buys almost nothing. In reality, slew/settle overheads, non-Lambertian scattering, correlated coronal systematics, and spatial cloud correlations impose diminishing returns. Furthermore, the distinction between *registered visits per pixel* and *number of physical spacecraft* must remain explicit: 64 registered visits can be delivered by 16 spacecraft operating over 4 successive 90-day campaigns, or by a 64-craft swarm in a single 90-day window.

7 Discussion and Limitations

Several simplifications bound this study. (i) We adopt the monopole, aperture-averaged SGL response; the solar quadrupole ($J_2 = 2 \times 10^{-7}$) creates astigmatic caustics that demand a J_2 -dependent deconvolution library (Turyshev & Toth 2021). (ii) Solar coronal background is modeled as Gaussian shot noise with mean subtracted, and Section 5.15 now measures rather than asserts the consequences of violating that. Coherent low-rank drift and streamer residuals are tolerated up to $\varepsilon \simeq 10^{-6}$ of the *coronal background*, which in planet-reference units is $\Delta y \simeq 0.08$ — comparable to the planet signal itself, since $Q_{\text{cor}}/Q_{\text{exo}} = 7.74 \times 10^4$ — and the tolerance knee sits at 3 – 4σ of single-dwell photon noise expressed in the same units ($1\sigma \simeq 3 \times 10^{-7}$). Calibration of the low-order spatio-temporal modes of the corona at the ppm-of-background level — not of the planet — is the requirement this places on the instrument, and a dedicated background differencing pass does not meet it: it costs a full photon- σ and removes only the mean (Turyshev 2026b). (iii) Clouds are modeled as 2D single-layer screens without 3D vertical structure, scattering phase functions, or specular ocean glint. The weather model is also homogeneous in the fiducial runs; Section 5.10 shows that the two structures which matter are surface coupling (a -0.198 paired loss that the cadence law cannot fix) and multi-timescale decorrelation (a marginal -0.025). The model is also not exhaustive in the surface it is tested against: Section 5.11 measures the morphology dependence at fixed land fraction and finds the pipeline insensitive to how finely the land is broken up but mildly sensitive — -0.040 in r , $\sim 12\%$ — to how coherently it is concentrated, because a single large continent generates a 10.5% larger cloud amplitude. A patchy high-albedo




Figure 8: Celestial placement (equatorial Mollweide) of the five confirmed targets (blue circles), candidate τ Cet e (green square), and their SGL focal directions (orange stars, purple triangle). The gray line marks the ecliptic plane. Ross 128 and Teegarden c focal lines lie in the ecliptic ($\beta_{\text{ecl}} \approx 0^\circ$), minimizing departure $\Delta\nu$.

cloud field, a heterogeneous climatology, or a venus-bright ocean world remain outside the tested set. (iv) Image-plane navigation is assumed known; uncompensated pointing jitter and ephemeris drift will act as an additional blur kernel.

(v) A limitation v2 did not acknowledge: the reconstruction grid. The surface is parameterized on 36×72 cells, whose spherical-harmonic support ends near $\ell \approx 18$ and whose solid-angle quadrature carries an orthonormality error of 0.018 at $\ell = 20$ (Section 5.3). Any statement about recovered resolution is therefore bounded by the grid, not by the optics, and the cloud-free $r(\ell) \geq 0.99$ result is a statement about the absence of rotational loss inside the supported range.

(vi) Statistical power. The headline ensembles are 10 paired seeds (6 for the campaign study), which sets a hard floor on paired significance of 1.95×10^{-3} and leaves several of the comparisons in this paper — the exact-versus-approximate deflation, the three benign climatologies, the arm A versus arm B difference — formally unresolved rather than resolved as equal. Where a difference is not resolved we now say so and give its confidence width instead of dropping the comparison. The physical cloud-only null has 40 members ($p \geq 0.024$ floor), and the profile-likelihood grid has 4 seeds, which is why its conclusion is a refutation (an exact invariance) rather than a fit.

(vii) Estimand discipline. Pearson r is exactly invariant to affine transforms of the retrieved map (Section 5.10), so it cannot support any statement about absolute albedo; SSIM and contrast are sensitive to different failure modes again. We report all three, and no single-metric claim about “reconstruction quality” should be read as an albedo claim.

(viii) Regularization selection. The one methodological

gap Section 5.13 exposes is that this paper has no truth-blind rule for the smoothing weight. The fiducial $\lambda = 3 \times 10^{-3}$ is frozen and defensible — the cloudy objective is flat to $\sim 0.1\%$ over five decades of it, so nothing in the data argues against it — but it is demonstrably not optimal: cloudy r keeps rising to $+0.178$ at $\lambda = 10$, and the control in Section 5.13 shows that gain is genuine $\ell \leq 4$ recovery rather than blur. A held-out campaign cannot resolve this, because correlation with a truth map is exactly what is unavailable in flight; a discrepancy-principle or GCV criterion, or a validated prior over surface models, is the missing piece and is stated as future work rather than implemented. The same section adds a constraint the opposite way: the assumed cloud amplitude σ_{cl} is a quantity the reconstruction barely uses (mean paired change in $r \leq 2.6 \times 10^{-4}$ over a $64\times$ range, and the same to 10^{-4} for the free profiler and v2 deflation) while χ^2 uses steeply, so it is a parameter that can be fit from residuals in flight without risking the map.

(ix) Coverage, not resolution, at high latitude. The raster supports $|\vartheta| \lesssim 85^\circ$, leaving the 144 exactly-zero basis columns of Section 4.4; Section 5.3 shows that roughly half of each polar latitude band is therefore set by the regularizer alone, and that the higher band correlation there is a low-contrast ice-cap artefact, not imaging. Polar structure is out of scope for this campaign design.

These limitations narrow, but do not remove, the core findings: rotation is tractable under characterized geometry and costs no measurable information inside the supported harmonic range; weather, not photon noise, sets the image quality by more than two orders of magnitude in variance; the cadence gain is an independent-looks effect whose size is set by the assumed weather decorrelation

time; and the algorithmic gain comes from eliminating the per-dwell-slot nuisance rather than from modeling the covariance of the weather.

8 Conclusions

1. Regularized generalized least squares on time-tagged samples (TDI), with exact anchored profiling of the per-dwell-slot cloud nuisance, recovers a rotating, cloud-free exoplanet with $r = 0.9786 \pm 0.0002$ (SSIM = 0.848 ± 0.001 , 10 paired seeds). Within the harmonic range the 36×72 grid supports ($\ell \leq 20$) no rotational band is lost; what remains is regularization and aperture. The v2 formulation “rotation introduces virtually zero information loss” is replaced by the measured statement: the phase-blind treatment of the same photons gives $r = 0.116$, so the time-registered operator is worth $+0.86$ in correlation, while temporal whitening itself is worth nothing at this cadence. A frozen-planet reference at the identical photon budget is not the other side of that comparison: it gives $r = 0.4074 \pm 0.0254$, *below* the rotating cloud-free result, because identical pose-passes make the time-domain design degenerate.

2. Clouds dominate the error budget. Per-sample cloud-induced variance exceeds coronal shot noise by 163–440 (12.8–21.0 in standard deviation) over 23.7–71.0% realized cover at $t_s = 1800$ s; fidelity collapses to $r = 0.249 \pm 0.010$ at the top of that range. The v2 “factors of 14–23” was a standard-deviation ratio produced by the rejected 3-node exposure quadrature and was described as though it were a variance ratio.

3. The cadence law is an independent-looks effect, and it is now measured under two declared resource ledgers. At fixed 8 h per raster position, $M_p : 4 \rightarrow 64$ takes r from 0.049 to 0.571 with every step unanimous across six paired seeds; forcing the same study into 89.87 d of wall clock changes the answer by less than the paired standard error at every M_p . What saturates is not r but N_{eff} : at $\tau_{\text{cloud}} = 4$ d, 64 revisits are worth 12.1 independent looks, so the cadence requirement must be quoted together with the assumed weather decorrelation time. Duty cycle falls from 99.4% to 90.5% over the tested range under the 45 s/dwell ledger; no configuration below $t_s = 429.6$ s was simulated, so the v2 claim about ultrafast overheads is deleted.

4. Ablating the estimator, not the pipeline, locates the algorithmic gain in nuisance elimination: exact profiling beats a fit that ignores the slot means by $\Delta r = +0.054 \pm 0.006$ (10/10 seeds, $p = 1.1 \times 10^{-5}$), while the difference between the exact treatment and v2’s approximate deflation, though material to the map (9.3% in norm), is not resolved by 10 seeds ($+0.003 \pm 0.002$, $p = 0.13$). The v2 claim that “slot deflation supplies the primary gain ($\Delta r = +0.041$, $p < 10^{-3}$)” attributed the effect to the approximation and quoted a p value below the 1.95×10^{-3} floor of a 10-seed sign test.

5. Spin-ephemeris tolerances are measured: 5° of pole

tilt costs 0.057 in r , 5° of assumed initial phase 0.047, 30° of phase leaves $r = 0.072$ and 60° leaves the map uncorrelated with the truth. We withdraw the claim that residual misalignments are refinable in flight via profile likelihood: the whitened χ^2 is numerically invariant ($< 10^{-11}$ relative) to assumed initial phase over $\pm 10^\circ$ and monotone in assumed period across the grid tested, because a phase shift is a re-labeling of a complete longitudinal basis. The spin state is observable through image quality, which peaks at truth on 4/4 seeds, but it is not recoverable from photometric residuals by grid search.

6. Resolution and detection are stated with their instrument limits. Cloud-free $r(\ell) \geq 0.99$ for all $\ell \leq 20$, the ceiling imposed by the reconstruction grid (quadrature orthonormality error 0.018 at $\ell = 20$); under 55.8% clouds a five-point running mean of $r(\ell)$ crosses 0.5 at $\ell \simeq 9$ (~ 2200 km at one Earth radius), so no per-seed $\ell_{\text{eff}} \approx 4$ –5 crossing is quoted. Land–ocean separation at that cover, $d' = 0.732 \pm 0.081$, exceeds all 40 members of the physical cloud-only null, giving $p \leq 0.024$ — the ensemble floor.

7. Mission mapping. Targets are a reproducible nominal-proxy table (Chen & Kipping radius relation with 4.0% dispersion on the terran branch, a $2.04 M_\oplus$ terran/neptunian transition, an inclination de-projection factor of 1.098, and an along-track velocity requirement of 0.11 – 5.8 km s $^{-1}$ per 90 d, Proxima Centauri b being the most demanding), not an uncertainty propagation; τ Cet e is the dynamically gentle case at $\sim 53 \times$ lower Δv . Science data volume is 2.10 MB per campaign and 10.4 kB per map, three orders of magnitude below the v2 claim, so downlink is not the binding constraint. Swarms remain an asymptotic architectural scenario.

8. Robustness is now measured rather than asserted, with two results that change how the pipeline should be used. Varying the weather parameters — correlation length 8° – 20° , advection 3° – 12° day $^{-1}$, decorrelation time 2–8 d — moves r by less than the seed-to-seed error, and the reason is that the reconstruction is nearly blind to the *assumed* cloud amplitude: over a $64 \times$ range of σ_{cl} the mean paired change in r is $\leq 2.6 \times 10^{-4}$ — for the anchored profiler, the free profiler and v2 deflation alike — while the whitened objective changes by a factor 8×10^4 . The pre-flight requirement is therefore on weather *structure*, and σ_{cl} is fittable in flight. Second, the smoothing weight is far from optimal and the objective cannot say so: cloudy r rises monotonically to $+0.178 \pm 0.007$ at $\lambda = 10$, a gain shown by a harmonic-band control to be genuine $\ell \leq 4$ recovery and not blur, while cloud-free r peaks at $\lambda = 0.3$ and turns over. The headline numbers keep the frozen weight, and the missing truth-blind selection criterion is stated as the open methodological problem. Surface morphology is bounded the same way: fragmenting the continents costs nothing measurable, concentrating them into one mass costs 0.040 in r .

Acknowledgments

This work used open-source software (NumPy, SciPy, Matplotlib, scikit-image). The author thanks the editor and referees for constructive reviews that substantially strengthened this analysis. No external funding was received.

References

- Anglada-Escudé, G., Amado, P. J., Barnes, J., et al. 2016, *Nature*, 536, 437
- Astudillo-Defru, N., Forveille, T., Bonfils, X., et al. 2017, *A&A*, 602, A88
- Bonfils, X., Astudillo-Defru, N., Díaz, R., et al. 2018, *A&A*, 613, A25
- Chen, J., & Kipping, D. 2017, *ApJ*, 834, 17
- Dreizler, S., Jeffers, S. V., Rodríguez, E., et al. 2020, *MNRAS*, 493, 536
- Dreizler, S., Luque, R., Ribas, I., et al. 2024, *A&A*, 684, A117
- Eshleman, V. R. 1979, *Science*, 205, 1133
- Feng, F., Tuomi, M., Jones, H. R. A., et al. 2017, *AJ*, 154, 135
- Friedman, L. D., Garber, D., Turyshev, S. G., et al. 2024, *Exp. Astron.*, 57, 1
- Helvajian, H., Rosenthal, A., Poklemba, J., et al. 2023, *J. Spacecr. Rockets*, 60, 829
- Madurowicz, A., & Macintosh, B. 2022, *ApJ*, 930, 19
- Aerospace Corporation & NASA NIAC 2022, Atomic Planar Power for Lightweight Exploration (APPLE), NIAC Phase II study
- Toth, V. T. 2025, *Res. Astron. Astrophys.*, in press (arXiv:2504.18630)
- Toth, V. T., & Turyshev, S. G. 2021, *Phys. Rev. D*, 103, 124038
- Toth, V. T., & Turyshev, S. G. 2023, *MNRAS*, 525, 5846
- Turyshev, S. G., & Toth, V. T. 2022c, *Phys. Rev. D*, 106, 044059 (arXiv:2207.06388)
- Turyshev, S. G., & Toth, V. T. 2023b, *Phys. Rev. D*, 107, 104063 (arXiv:2302.08102)
- Turyshev, S. G. 2026a, *Phys. Rev. D*, 113, 023034 (arXiv:2506.20236)
- Turyshev, S. G. 2026b, arXiv:2606.14899
- Turyshev, S. G. 2026c, arXiv:2606.18300
- Turyshev, S. G. 2026d, arXiv:2602.04198
- Turyshev, S. G. 2026e, arXiv:2606.29138
- Turyshev, S. G., & Toth, V. T. 2017, *Phys. Rev. D*, 96, 024008
- Turyshev, S. G., & Toth, V. T. 2019, *Phys. Rev. D*, 99, 024044
- Turyshev, S. G., & Toth, V. T. 2020, *Phys. Rev. D*, 102, 024038
- Turyshev, S. G., & Toth, V. T. 2021, *Phys. Rev. D*, 103, 064076
- Turyshev, S. G., & Toth, V. T. 2022a, *MNRAS*, 515, 6122
- Turyshev, S. G., & Toth, V. T. 2022b, *Phys. Rev. D*, 105, 044012
- Turyshev, S. G., Shao, M., Toth, V. T., et al. 2020, Direct Multipixel Imaging and Spectroscopy of an Exoplanet with a Solar Gravity Lens Mission, NASA NIAC Phase II Final Report